

MARKET EVIDENCE BRIEF

Brand visibility inside AI assistants: is there demand, and does it work?

Whether businesses pay to be recommended inside ChatGPT, Gemini, Perplexity, Copilot and Claude, by buyer segment, and what evidence shows it pays off

Research only: evidence, not a verdict

Evidence as of 2026-09

Some businesses already spend and one AI engine sells ads at a \$1 billion run rate, but no published case proves a sales effect

8 of 27

buyer segments show a spend signal (strict reading[†])

Segment = kind of placement × industry × buyer size. Analyst-derived, 2026-09-23

\$1 billion

ChatGPT ads: "\$1 billion in annualized revenue run rate"

Company-stated, 2026-08-31

None

of ~980 cases screened to 2026-09-22 and ~1,640 more on 2026-09-23 used a holdout, geo-split or switchback test

Measured-by-us, 2026-09

Where the evidence stops: the market **exists** (proven); its **effect** rests on vendors' own reports; its **size** is assumed. No kind of placement has a measured market size.

[†] A second reading of the same figure exists; both readings are in the appendix.

A brand can appear inside an AI assistant's answer three ways; each is a separate market

Organic recommendation

Unpaid appearance in an answer.
Sold as GEO (generative engine optimisation), a cousin of SEO (search engine optimisation).

Paid placement

Ads shown inside or beside the AI answer, bought from the engine.

Agentic commerce

The assistant's agent picks the product and executes the checkout for the user.

The question. Does real demand exist, per buyer segment, and what evidence shows it works?

Buyer segment. Kind of placement × industry × buyer size gives 27 segments ("cells"). Industries tracked: skincare and beauty; B2B SaaS (business-to-business software as a service); high cost-per-acquisition regulated.

Three metrics, never mixed:

- **Visibility:** the brand appears in the answer
- **Traffic:** sessions referred to the brand's site
- **Sales:** revenue attributed to the answer

A visibility gain is not evidence of a sales gain.

Every figure carries a source kind, a tier and a date; tier 1 is strongest, 7 weakest

Tier	Provenance
1	Measured by us
2	Filed: SEC (US Securities and Exchange Commission) forms, court exhibit, funding filing
3	Platform primary: own docs, pricing page
4	Panel or clickstream data
5	Vendor or agency study with sample size, dates, method
6	Vendor blog, no sample size: not a source
7	Listicle: not pulled

Source kinds: filed · company-stated · vendor-reported · analyst-derived (forecasts, our compilations) · measured-by-us.

Case grades for "did it work" stories:

- **Gold:** holdout, geo-split or switchback test
- **Silver:** before/after with an unaffected control
- **Bronze:** visibility change only

Strict rule: any of the seven evidence items missing on the case's own page caps it at Bronze.

Strict vs loose segment reading: strict counts a segment only if every signal was run; loose counts a blocked, logged channel as checked.

From 2026-02 to 2026-09 engines opened ad inventory, publishers filed traffic losses, and regulators named AI search

Engines opened ad inventory	Publishers filed losses	Regulators and capital
2026-02-04 Anthropic: "<i>Claude will remain ad-free</i>" 2026-02-09 ChatGPT ads launch in the US 2026-02-18 Perplexity ads reported "<i>winding down</i>"† 2026-08-18 ChatGPT ads in 31 European markets 2026-08-31 ChatGPT ads in "<i>over 40 countries</i>"	2026-02-25 NerdWallet: credit-cards revenue "<i>decreased 24% year-over-year</i>", citing "<i>AI overviews and LLMs</i>" (large language models) 2026-05-04 IAC / People Inc: "<i>63% decline in Google Search referrals over two years</i>" 2026-08-03 Filing: "<i>22% decline in Core Sessions</i>" 2026-09-17 Microsoft court exhibit: "<i>83-93% drops in click-through rates</i>"	2026-03-02 Semrush: "<i>AI products surpassed \$38 million in ARR</i>" (annual recurring revenue) 2026-04-28 Semrush to be acquired by Adobe 2026-06-17 UK CMA (Competition and Markets Authority) fair ranking covers "<i>search generative AI features</i>" 2026-08-31 EU designates ChatGPT a VLOSE (Very Large Online Search Engine); duties from January 2027 2026-09-01 Yext buys GoShine 2026-09-15 Profound: "<i>\$180M Series D at \$1.8B valuation</i>"†


No kind of placement has a measured market size; only floors built from disclosed revenue exist

Kind of placement	Floor: disclosed revenue only	Source kind, tier, date
Organic recommendation	"\$42.2M–\$48.2M + €2.2M ARR", 4 vendors†	filed (2) and company-stated (5), 2025-11 to 2026-05
Paid placement	"≥ \$1B run rate", one engine (ChatGPT)	company-stated (3), 2026-08-31
Agentic commerce	unknown — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon	measured-by-us, 2026-09-23

Floors are not market sizes. They add up only what named companies disclose. The paid floor is one engine's run rate.

Price paid is not published. *"0 of 8 publish a rate card"* among AI engines (company-stated, tier 3, 2026-09-22).

Every published market size is a forecast, and forecasts for the same market disagree by up to 26.3×



Bars span the lowest to highest published forecast. Horizons differ (2034 vs 2030); compare spreads, not totals. A second computation of the agentic spread reads "roughly 35×" over 2029–2030†.

GROWTH

ChatGPT passed 1 billion weekly users and AI-referred site visits rose 117.4% year on year

ChatGPT weekly active users

700M → ">1 billion"

2025-09-15 → 2026-08-31

Company-stated, tier 3

Visits to websites referred by AI assistants


**770.7M / month,
"+117.4%"**

year on year, June 2025 – May 2026

Vendor-reported, tier 4 (panel / clickstream)

Usage and referral traffic are growing. Neither figure measures what brands spend or what sales they gain; those are the next two slides.

Spend shows in 8 of 27 buyer segments, mostly organic; willingness to pay is unknown in all 27



One square = one segment; squares show counts, not segment positions. Blank = not checked.

Spend segments by kind (strict): organic 6+, paid 1, agentic 1. **Small business:** one spend segment only, B2B SaaS organic.

No filing names the budget this spend comes from; the free substitute is doing nothing

Budget line: unknown — checked all three industries, 2026-09-23. The nearest filed statements (filed, tier 2):


Coty, 8-K current report, 2026-08-19 (beauty): *"rightsizing ... global brand marketing functions" beside "optimizing the visibility and recommendation of our brands across AI platforms"*

eHealth, 10-K annual report FY2025 (high cost-per-acquisition regulated): SEO and paid search *"may be adversely impacted"; "could further increase our marketing expenditures"*

Semrush, 10-K annual report FY2025 (seller's own framing): customers *"supplement their traditional SEO tools with newer GEO and AI search products"*

Substitute: do nothing. Google states *"no additional requirements... nor other special optimizations necessary"* to appear in its AI features (company-stated, tier 3, 2026-09-22).

Of 34 visibility vendors, none discloses the prompt set behind its score, and 2 file revenue



Self-serve prices run from "\$20/mo" to "\$999/month" (vendor-reported, tier 3, 2026-09-22).

Largest filed AI line: Semrush "*AI products surpassed \$38 million in ARR*" (filed, tier 2, 2026-03-02). 43 company profiles compiled (measured-by-us, 2026-09-23).

Engines sell the ads, the measurement and the checkout themselves; gaps remain as hypotheses

Engines keep the rails

- OpenAI ships an Ads Manager, a Pixel, a Conversions API (application programming interface) and "*more than 50 technology and measurement partners*" (company-stated, tier 3, 2026-08-31).
- Agent checkout: of 4 live programmes, 1 states a fee. Copilot "*does not take a commission*" (company-stated, tier 3, 2026-09-22).
- Payment-protocol fees: 6 of 6 clauses unknown — checked, 2026-09-22.

Gaps: hypotheses, not plans


- A visibility score with a disclosed prompt set.
- Independent measurement of ads inside answers. Adthena lists "\$350/mo" and "\$399/mo" (vendor-reported, tier 3, 2026-09-23)[†].
- Compliance tooling for EU ad-repository duties: unknown — checked, 2026-09-23.
- Brand-side confirmation of vendor claims: 1 of 168 brands[†] (measured-by-us, 2026-09-23).

That the market exists is proven; that it works rests on vendors' own reports; its size is assumed

Rung	What the record holds	Kind, tier
Proven	Four† engines sell ads inside answers (2026-09-23); publisher referral declines are filed	company-stated 3; filed 2
Inferred, positive	0 Gold; Silver 7 as graded, 1 under strict rule (2026-09-22); 4 more, 1 under strict rule (2026-09-23)	vendor-reported 5
Inferred, negative or null	2 of 7 Silvers (2026-09-22), 4 more graded (2026-09-23); 0 of 15 advertiser ad results controlled	filed 2 to vendor-reported 5
Assumed	Every published market size is a forecast	analyst-derived 5–6
Own measurement	0 ad units in 76 Google AI Mode runs, Vietnam network, 2026-09-22	measured-by-us 1

One brand confirms a vendor result itself: Chime, *"tripled our AI citations"* (company-stated, tier 3, 2025-11-24). AI Mode = Google's conversational search tab.

Of 52 graded cases, 10 or 5 are experimental depending on the reading; 13 report an action only



The two readings differ on whether a before/after comparison with a control counts as experimental or observational. Neither reading yields a Gold-grade case. Cases are vendor-reported (tier 5), 2026-09-23.

Referral collapse and platform capture each touch all 27 buyer segments

Risk	Strongest evidence (kind, tier, date)	Early signal
Referral collapse	"22% decline in Core Sessions" (filed, 2, 2026-08-03)	filings naming "AI Overviews"
Regulation	ChatGPT VLOSE 2026-08-31; ad-repository duty January 2027 (filed, 2)	EU designation list
Vendor consolidation	Semrush into Adobe 2026-04-28; Yext buys GoShine 2026-09-01 (filed, 2)	acquisition filings
Platform capture	Engine ads managers, pixels, zero-commission checkout (company-stated, 3, 2026-09)	engine help centres
No payback evidence	Vendor's own campaigns: CTR 1.30%, "very few sign-ups" (vendor-reported, 5, 2026-08-10)	vendor experiment pages
Measurement opacity	Google Search & Other "over \$63 billion", no AI line (company-stated, 3, Q2 2026)	earnings releases
Substitute suffices	Optimisation levers move citation on "none" of ten engine families (analyst-derived, 4, 2026-09-07)	replications
Engine policy reversal	Claude ad-free; Perplexity ads "winding down" (analyst-derived†, 5, 2026-02-18)	engine newsrooms

Ranked by strongest evidence tier, then breadth. Likelihood, impact, mitigant: no source states one. CTR = click-through rate.


WHAT IS MISSING

Price paid, market size and causal sales effect are unpublished; some small-business demand sits behind bot walls

Missing	Why	Cost to close, per the record
Measured size, any kind of placement	analyst reports paywalled	Forrester "individual purchase (\$1495)" (company-stated, 2026-09-23)
Price paid; engine rate cards	not published: "0 of 8 publish a rate card" (company-stated, 3, 2026-09-22)	no source states one
Causal sales lift	no holdout, geo-split or switchback published	no source states one
Job-board and EU tender demand (Indeed; TED, Tenders Electronic Daily)	bot walls	"a human browser session would clear them"

HOW STRONG IS THE EVIDENCE OVERALL

44.0% of load-bearing claims reach tier 3 or better, against an 80% bar



The 25-claim main list scored 68.0% under the tiers assigned 2026-09-22 and 44.0% after re-tiering on 2026-09-23. Both readings fall short of the bar.

No decision is asked; the evidence stops at existence, and three readings are yours to choose

1. Which market

Organic, paid, agentic, or several. Organic has the most spend segments; paid has the one large disclosed revenue; agentic has forecasts 26.3× apart and no floor.

2. Which proof reading

Silver cases as the source graded them, or under the strict rule that caps incomplete cases at Bronze.

3. Which demand reading

Strict or loose. Both show 8 spend segments; strict leaves 1 segment blank, loose 0.

The research produces evidence, not a verdict. Whether to act on it is the director's call.

Terms and acronyms

Term	Meaning
AI	Artificial intelligence
AI Mode; AI Overviews	Google's conversational search tab; Google's AI answer block on a results page
API	Application programming interface
ARR	Annual recurring revenue
B2B SaaS	Business-to-business software as a service, a tracked industry
CMA	UK Competition and Markets Authority
CTR	Click-through rate
DSA	EU Digital Services Act, the law behind the VLOSE designation
EU; UK; US	European Union; United Kingdom; United States
FY; Q2	Fiscal year; second calendar quarter
GEO	Generative engine optimisation, the organic market's trade name

Terms and acronyms

Term	Meaning
Gold / Silver / Bronze	Case grades: holdout, geo-split or switchback / before and after with an unaffected control / visibility change only
LLMs	Large language models
Experimental / observational reading	Two readings of "before/after with control" when counting study designs
SEC; 8-K, 10-K	US Securities and Exchange Commission; its current-report and annual-report forms
SEO	Search engine optimisation
Strict rule	Any of the seven evidence items missing on the case's own page caps it at Bronze
Strict / loose segment reading	Strict: every signal must be run; loose: a blocked, logged channel counts as checked
TED	Tenders Electronic Daily, EU public procurement notices
VLOSE	Very Large Online Search Engine, a DSA designation
†	A second reading or date of the figure exists; see Appendix B

Where the record holds two readings or two dates

Figure in the deck	Other reading or date
Segments, strict: 8 spend / 1 attention / 17 none / 1 blank (2026-09-23)	Loose 8 / 1 / 18 / 0 (2026-09-23); 7 / 1 / 11 / 8 (2026-09-22), six of seven spend reads on job postings
Organic floor "\$42.2M–\$48.2M + €2.2M ARR"	"\$5.3M–\$35.2M annualised", price × count build, 4 of 34 rostered
Agentic spread 26.3× (2030)	"roughly 35×" (2029–2030 horizon)
Spend segments organic 6 / paid 1 / agentic 1 (2026-09-23)	Organic 5 / paid 1 / agentic 1 (2026-09-22)
Profound round labelled company-stated, tier 3	Labelled vendor-reported, tier 3
Brand corroboration 1 of 168 (cumulative)	0 of 59 (2026-09-22)

Where the record holds two readings or two dates

Figure in the deck	Other reading or date
Four engines sell ads (2026-09-23)	"5 of 8 engines carry a live product" (2026-09-22)
Experimental reading 10 / 29 / 13 of 52	Observational reading 5 / 34 / 13 of 52
Perplexity wind-down labelled analyst-derived, tier 5	Labelled company-stated, tier 5
Share at tier 3 or better: 11 of 25 = 44.0% (2026-09-23)	17 of 25 = 68.0% (2026-09-22 tiers)
Ads-in-answers measurement: Adthena "\$350/mo", "\$399/mo"	Sensor Tower: "Contact Sales" only

What this deck leaves out, and how it was built

- **Omitted for length; held in the source research:** manipulation-technique research, publisher licensing, per-country EU figures, funding rounds beyond Profound, the full event timeline and the full risk register.
- **Method:** visibility, traffic and sales are never mixed. Each success case carries its grade as its source graded it and its grade under the strict rule.
- **Numbers:** every figure is carried as the source states it: no rounding, no averaging, no new arithmetic. Where two figures conflict, both appear.
- **Charts:** plot only figures stated in the record. Bar positions on the forecast chart use a logarithmic scale; labels carry the stated figures verbatim.
- **Internet sources only:** no interviews or outreach.