

Do brands really pay to be recommended inside AI assistants, and does it work?

Brand visibility and product recommendation inside ChatGPT, Gemini, Perplexity, Copilot and Claude: demand by segment, and the evidence that it pays.

Research only: evidence, not a verdict

Evidence as of 2026-09

Every figure carries its source kind and date. The decision belongs to the reader.

Demand is thin and enterprise-heavy, ad inventory is real, and no published case shows a sales effect

8 of 27

buyer segments show money moving;
willingness to pay unknown in all 27
measured-by-us, 2026-09-23

\$1 billion

"annualized revenue run rate" for
ChatGPT ads, live since 2026-02-09
company-stated, 2026-08-31

0 Gold

controlled sales-effect cases in ~980
plus ~1,640 more candidates screened
measured-by-us, 2026-09-22 and 2026-09-23

Three sub-markets, never merged: each has its own size, buyers and evidence

Sub-market	What the brand gets	Who sells it
Organic	Named in an AI answer, unpaid; tools to track and improve that	Visibility-tool vendors, agencies
Paid	Ads shown inside AI answers	The assistant owners (engines)
Agentic	The assistant completes the purchase for the shopper	Engines, payment and commerce platforms

Each figure is scored by where it came from; lower tier is stronger

Tier	Source	Label used here
1	Measured by this programme	measured-by-us
2	Filed: annual reports, court exhibits, regulator decisions	filed
3	Platform's own docs, pricing or newsroom page	company-stated
4	Panel, clickstream or telemetry data; published replication	vendor-reported; analyst-derived
5	Vendor or agency study with sample, dates, method	vendor-reported
6	Vendor blog without sample; weakest forecasts	vendor-reported; analyst-derived
7	Listicle or roundup: not used	none

Case grades for "does it work"


- **Gold:** holdout, geographic split or switchback test
- **Silver:** before and after, with an unaffected control
- **Bronze:** visibility changed; sales not shown
- **Fools gold:** revenue claim with no baseline or control

Under the strict reading of the grading bar, a case missing any one of seven required items is Bronze at best.

Only 8 of 27 buyer segments show money moving, under both readings

A segment is one sub-market × one vertical (skincare and beauty; business-to-business (B2B) software; high cost-per-acquisition regulated) × one buyer size. Each is read as **spend** (money moved), **attention** (interest only), **none** (checked, empty) or **blank** (not checked).

Segments, count (of 27), measured-by-us, 2026-09-23




Willingness to pay: unknown in all 27 segments. No buyer has disclosed a price paid.

An earlier tally, dated 2026-09-22, exists; both are in the appendix.

Where money shows, it is mostly organic and enterprise, and much of it is hiring

Segments reading spend, count by sub-market, 2026-09-23



The only small-business spend segment: B2B software, organic.

- Six of the seven spend reads of 2026-09-22 rested on job postings: *"headcount budget, not category spend"* (measured-by-us).
- Enterprises hiring for it: e.l.f., Pennylane, Cigna postings (company-stated, tier 3). Intent, not spend.
- Outside the 27 segments:
 - Local multi-location businesses: 9 segments, spend in 1
 - European Union (EU), per country: 45 segments, spend in 2

DOES IT WORK?

No published case proves a sales effect; the best cases show visibility, not revenue

0 Gold

No holdout, geographic split or switchback in ~980 candidates screened to 2026-09-22, or ~1,640 more on 2026-09-23
measured-by-us

7 vs 1

Silver cases as graded vs under the strict seven-item bar (2026-09-22); a later screen added 4 as graded, 1 strict (2026-09-23)
measured-by-us; cases vendor-reported, tier 5


1 of 168

brands named in vendor case studies that confirm the case themselves (Chime)
measured-by-us, 2026-09-23

Negative evidence sits beside the positives

- 2 of 7 first-round Silver cases were negative or null, including NerdWallet (filed, tier 2).
- Paid ads: 0 of 15 reported ad results had a control (measured-by-us, 2026-09-23).
- A replication found 2023 optimisation levers move citation on "none" of ten engine families (analyst-derived, tier 4, 2026-09-07).

2026 put ad inventory, new rules and vendor deals on the calendar



Time axis to scale, 2026-02 (left) to 2027-01 (right). VLOSE: Very Large Online Search Engine, an EU Digital Services Act designation; ChatGPT has 159.1M EU recipients and must keep an ad repository. UK regulator: Competition and Markets Authority (CMA).

AI answers are taking clicks away from publishers, and filings say so

"decreased 24%"

NerdWallet credit-cards revenue
filed, 8-K current report, 2026-02-25

"63% decline"

People Inc: "in Google Search referrals
over two years"
filed, investor deck, 2026-05-04

"83-93% drops"

Microsoft: "in click-through rates"
filed, court exhibit, 2026-09-17


The three figures measure different things: revenue, referral traffic, click-through rate. They are shown side by side, never combined.

This is the strongest evidence in the programme (filed, tier 2). It shows the old channel losing value. It does not show the new channel paying back.

HOW BIG: AUDIENCE

The audience is large and still growing

ChatGPT weekly active users, millions
company-stated, tier 3



Bars to scale from zero. The 2026 bar is drawn at the stated floor; the source says only "> 1 billion".

770.7M / month

visits referred by AI assistants, "+117.4%" year on year, Jun 2025 to May 2026
vendor-reported, tier 4

\$224,532M

Google Search & other revenue: the classical search baseline AI answers sit beside
filed, tier 2, fiscal year 2025


No one has measured the market; the floors that can be proven are small

Sub-market	Floor from disclosed revenue	Source kind, date
Organic	\$42.2M–\$48.2M + €2.2M annual recurring revenue (ARR), 4 vendors. Another basis gives a different floor (appendix).	filed (tier 2) and company-stated (tier 5), 2025-11 to 2026-05
Paid	≥ \$1B run rate, ChatGPT only	company-stated, tier 3, 2026-08-31
Agentic	<i>unknown — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon 2026-09-23</i>	no figure disclosed

A floor is what disclosed revenue proves exists. It is not a market size. Every published market size is a forecast: 18 figures, 15 forecasts, 0 measured (analyst-derived, tier 6).

Published forecasts disagree by up to 26.3 times within one sub-market

Highest forecast ÷ lowest forecast, same sub-market and year, times (x)



Range	Tier
Organic 2034: \$17.15B–\$32.92B	6
Paid 2030: ~\$5B–\$100B+, scopes differ	5
Agentic 2030: \$190B–\$5T	5–6

Dashed bar: an earlier count of the same forecasts, "roughly 35x" over 2029–2030 (2026-09-22). Both stand.

Which budget this spend comes from is unknown; three filings point near it


Budget line it draws from: unknown — checked US filings 2026-09-23.

Company, filing	What it says	Kind, date
Coty, 8-K	Pairs brand-marketing cuts with " <i>optimizing the visibility and recommendation of our brands across AI platforms</i> "	filed, 2026-08-19
eHealth, 10-K annual report	Names search engine optimisation (SEO) and paid search as at risk	filed, fiscal year 2025
HubSpot	Bought an answer-engine optimisation platform for "\$16.5 million"	filed, 2025-12-01

These show where the money may come from: search and brand-marketing budgets. None states a line item for AI-assistant visibility.

34 organic-visibility vendors, funded and consolidating, but few disclose what they sell

Vendors disclosing each item, count of 34, measured-by-us, 2026-09



Roster: 13 pure-plays, 16 incumbents bundling a feature, 5 agencies. A "fixed prompt set" is the published list of questions a tool uses to measure visibility.

- **Money in:** Profound "\$180M Series D at \$1.8B valuation", 2026-09-15 (company-stated, 3).
- **Consolidation:** Semrush into Adobe, "\$12.00" a share, 2026-04-28; Yext bought GoShine, 2026-09-01 (filed, 2).
- **Conflict:** Scrunch into Sitecore, 2026-06-03: "not disclosed" (company-stated, 3) vs "\$225 Million" (press, 5).
- **Where a hypothetical product could sit:** open; no evidence addresses it.

The assistant owners run the paid auction and the checkout themselves

Paid: who sells ads inside answers

Engine	Status
ChatGPT	Self-serve; "\$3–\$5 USD per click" recommended max bid
Google	Existing campaigns auto-eligible
Amazon Rufus	Live 2026-03-25
Claude	"will remain ad-free" (2026-02-04)
Perplexity	"winding down... by the end of 2026" (press, 5)

Four engines sell ads (2026-09-23); an earlier count on another basis exists.

Agentic: who runs checkout

- 4 engines run checkout: ChatGPT, Google, Copilot, Perplexity.
- Checkout protocols (ACP, UCP, x402) are open source, Apache-2.0 licence.
- Copilot Checkout "does not take a commission or affiliate fee".
- Fee clauses: *unknown in 6 of 6 protocols* (2026-09-22).

Engines tell brands little


- OpenAI and Anthropic ship 0 native visibility tools (measured-by-us).
- Google: *"no additional requirements... nor other special optimizations necessary"* (2026-09-22).

Existence is proven, demand is inferred, payback is assumed

Rung	Claim	Best evidence
Proven	Paid inventory exists at scale	Four engines sell ads, 2026-09-23; \$1B run rate (company-stated, 3)
Proven	AI answers cut publisher referrals	8-K, investor deck, court exhibit (filed, 2)
Inferred	Enterprises hire for it	e.l.f., Pennylane, Cigna postings (company-stated, 3); intent, not spend
Inferred	AI referrals convert better	Adobe: "Is Now 54% Higher", May 2026 (vendor-reported, 4); earlier reading exists
Inferred	Content changes move visibility	7 Silver as graded, 1 strict (2026-09-22); 4 and 1 more (2026-09-23); vendor-reported, 5
Assumed	Any intervention lifts sales	0 Gold; 0 of 15 ad results controlled (measured-by-us, 2026-09-23)

Most proof is observational; the experimental count depends on one reading

Cases by evidence class, count, measured-by-us, 2026-09-23. Scale: 6 pt per case.
"Control as experimental" and "control as observational": two readings of before-and-after-with-control cases.



Experimental: a controlled test. Observational: before and after, no clean control. Action only: the move is named, the outcome is not. The 4 dashed cases sit outside the classes.

Referral loss and regulation carry the strongest evidence of any risk

#	Risk	Strongest evidence (tier)	Early signal
1	Referral collapse	"83-93% drops" (2); "decreased 24%" (2)	Filings naming "AI Overviews"
2	Regulation	EU VLOSE duty January 2027 (2); UK CMA (2); Google AI-answer ads exclude finance, healthcare (3)	EU designation list
3	Vendor consolidation	Semrush–Adobe (2); Yext–GoShine (2)	Current reports on rostered vendors
4	Platform capture	OpenAI Ads Manager, Pixel, "more than 50 technology and measurement partners" (3)	Engine help centres
5	No payback evidence	1.30% click-through, "very few sign-ups"; "just 3%" of \$250K spent (5)	Filings naming "ChatGPT ads"
6	Measurement opacity	ChatGPT ad presence 0.8%–26% across panels, unreconciled (5)	Vendor method pages
7–8	"Do nothing" suffices; engines reverse policy	Google guidance (3); Claude ad-free (3)	Google guidance page; engine newsrooms
9–11	Agentic fees; manipulation; unsizeable market	Spam policy on "manipulate generative AI responses", 2026-08-28 (3); 18 figures, 0 measured (6)	Protocol repos; spam-policy page; forecast sales pages

Ranked by evidence tier, then breadth (2026-09-23), not by severity. Likelihood, impact and mitigation: no source states one.

The decisive numbers are missing: no price paid, no disclosed method, key channels blocked

Missing	Why, and cost to close
A price any buyer paid	<i>none disclosed; unknown — not in repo</i>
Engine rate card (0 of 8); protocol take rate	<i>none disclosed; unknown — not in repo</i>
Analyst size of each sub-market	Paywalled; Forrester "individual purchase (\$1495)"; others <i>unknown — not in repo</i>
Vendor funding, 17 of 60 rows	Crunchbase paywall; <i>unknown — not in repo</i>
Small-business hiring	Indeed bot wall; "a slower human session"

21 of 31

channels probed were reached
measured-by-us, 2026-09-23

Blocked or paywalled: Indeed, EMARKETER, live Reddit, Crunchbase, Gartner, Forrester. The gaps are blocked, not exhausted.

No decision is asked; two readings of the evidence are yours to set

1. Which case count governs

Controlled cases that content changes move visibility;
cases vendor-reported, tier 5.

As graded: 7 Silver (2026-09-22), plus 4 (2026-09-23).

Strict seven-item bar: 1 Silver, plus 1.

2. Which segment reading governs

Strict: 26 of 27 segments closed (checked).

Loose: 27 of 27 closed.

Both readings give 8 spend segments (measured-by-us,
2026-09-23).

What the evidence supports: inventory exists, the old search channel is losing clicks, enterprises are paying attention.

Where it stops: no price paid, no measured market, no controlled sales effect.

Terms and acronyms

Term	Meaning
ACP	Agentic Commerce Protocol: OpenAI and Stripe checkout specification
AEO / GEO	Answer engine / generative engine optimisation: names for the organic sub-market
AI; AI Overviews	Artificial intelligence; Google's AI answer block on a search results page
ARR	Annual recurring revenue
B2B	Business to business
CMA	UK Competition and Markets Authority
CPA	Cost per acquisition
CPC / CPM / oCPC / oCPM	Cost per click / per thousand impressions / their optimised-bidding variants
CTR	Click-through rate: clicks divided by impressions
DSA	EU Digital Services Act
EU / UK / US	European Union / United Kingdom / United States

Terms and acronyms, continued

Term	Meaning
FY	Fiscal year
M / B / T; K	Million / billion / trillion; thousand
Rufus	Amazon's shopping assistant
SEC; 8-K / 10-K / 10-Q	US Securities and Exchange Commission; its current, annual and quarterly report forms
SEO	Search engine optimisation, classical
Spend / attention / none / blank	Segment reads: money moved / interest only / checked and empty / not checked
Strict / loose	Strict: every signal must be run. Loose: a logged blocked channel counts as checked
Tier	Provenance score, 1 (measured by us) to 7 (listicle); lower is stronger
UCP / x402	Universal Commerce Protocol (Google-founded) / a payment protocol built on HTTP status 402
USD / EUR	US dollar / euro
VLOSE	Very Large Online Search Engine: DSA designation; obligations include an ad repository

Where the evidence gives two figures, both stand; none is averaged

Figure A (kind, date)	Figure B (kind, date)	Why two
Organic floor "\$5.3M–\$35.2M annualised", price × count, 4 of 34 vendors; vendor-reported inputs, tier 5, 2026-09-22	"\$42.2M–\$48.2M + €2.2M" ARR; filed 2 / company-stated 5, 2026-09-23	two bases
Segment tally: spend 7 / attention 1 / none 11 / blank 8; measured-by-us, 2026-09-22	Strict 8 / 1 / 17 / 1; loose 8 / 1 / 18 / 0; measured-by-us, 2026-09-23	two dates, two readings
Agentic spread "roughly 35×", 2029–2030; analyst-derived, 2026-09-22	26.3×, 2030 only, 7 forecasts; analyst-derived 5–6, 2026-09-23	two bases
Silver 7 as graded / 1 strict; ~980 screened; measured-by-us, 2026-09-22	+4 as graded / +1 strict; ~1,640 more screened; measured-by-us, 2026-09-23	two readings, two dates
Brand confirmation 0 of 59 checked, ~107 not checked; 2026-09-22	1 of 109 (Chime); cumulative 1 of 168; 2026-09-23	two dates

Two figures, both attributed, continued

Figure A (kind, date)	Figure B (kind, date)	Why two
Spend reads resting on postings: six of the seven; measured-by-us, 2026-09-22	"headcount budget, not category spend", no restated share; 2026-09-23	two dates
Adobe AI-referral conversion "42% Higher", Mar 2026; vendor-reported 4	"54% Higher", May 2026; same source	two dates
Shopify "nearly 50% higher rates than organic search", Q1 2026; vendor-reported 5	"converted about 80% better", Q2 2026	two dates
Peec AI ARR "more than \$4 million", 2025-11-17; company-stated 5	"10m", 2026-05; company-stated 5	two dates
Scrunch acquisition "not disclosed"; company-stated 3, 2026-06-03	"\$225 Million"; analyst-derived 5, 2026-06-03	two tiers

Two figures, both attributed, continued

Figure A (kind, date)	Figure B (kind, date)	Why two
ChatGPT ad presence ~0.8% of 500+ prompts, launch weeks; vendor-reported 5	4.47% of ~850K US, Mar–May; 26% US desktop, Jun; vendor-reported 5	several panels
Skincare × agentic × enterprise reads spend; measured-by-us, 2026-09-22	Reads attention; 2026-09-22	two compilers
Engines with a live paid product "5 of 8"; company-stated 3, 2026-09-22	"Four engines sell ads inside AI answers"; company-stated 3, 2026-09-23	two counting bases
OpenAI Instant Checkout fee: unknown — checked 2026-09-22	"Merchants pay a small fee on completed purchases", rate undisclosed; 2026-09-23	two dates
Tier-3 share of load-bearing claims "17 of 25 = 68.0%"; 2026-09-22	"12 of 27 = 44.4%"; "11 of 25 = 44.0%"; 2026-09-23	two lists, two tier bases

Published prices exist; none is a price a buyer paid

Seller	Price, verbatim	Kind, tier, date
Organic tools, 13 pure-plays	"\$20–\$999/month self-serve; Enterprise 'Custom' at nearly all"; disclosed 14 of 34	vendor-reported 3, 2026-09
Semrush	"\$199/mo Starter"; "SEO base \$139/mo"; delta "+\$60/mo monthly"	vendor-reported 3, 2026-09-22
Pace Generative (agency)	"\$1,499 / \$1,999 / \$2,499" per package; "4 contact-only" agencies	vendor-reported 3, 2026-09-22
OpenAI, ChatGPT Ads	"\$3–\$5 USD per click" max bid; minimum daily spend "25 USD"; "CPM, CPC, oCPC, oCPM"	company-stated 3, 2026-09-23
OpenAI, trade-reported	Launch "\$60" CPM, "\$200,000–\$250,000 minimum"; "\$25" CPM mid-April; minimum removed 2026-05-05	analyst-derived 5, 2026-02 to 2026-05
Google, Microsoft, Amazon	Auction, no unit price; existing campaigns eligible	company-stated 3, 2026-09-22
Kontext (ad network)	"Rates starting from just \$3 CPM"	vendor-reported 3, 2026-09-23
Shopware (commerce platform)	Rise "From € 600/month"; Evolve "€2,400/mo"	vendor-reported 3, 2026-09-22

Where the figures come from, and how they were handled

Source	Kind, tier
SEC filings, court exhibits, EU and UK regulator decisions	filed, 2
Engine help and newsroom pages; job postings	company-stated, 3
Vendor pricing and newsroom pages	vendor-reported, 3
Similarweb, StatCounter, Adobe Analytics panels; academic replication	vendor-reported / analyst-derived, 4
Vendor case studies; ad-intelligence panels; Reddit archive	vendor-reported, 5
Market forecasts (Dimension, Market Decipher, EMARKETER, WPP, Morgan Stanley, McKinsey)	analyst-derived, 5–6
Own screen counts, segment tallies, channel probes	measured-by-us, 1

- Internet sources only; no interviews.
- No figure recomputed, rounded or converted; dollars and euros carried as stated.
- Where one figure has two readings or dates, the main slides carry the later one; Appendix B carries both.
- Risks ranked by evidence tier, then breadth, not severity.