

Brand visibility inside AI assistants: demand shows in job postings and ad inventory, not yet in payback

What the published evidence says about who buys, how big the market is, who sells, and whether it works

Research only: evidence, not a verdict

Evidence as of 2026-09

AI assistants covered: ChatGPT, Gemini, Perplexity, Copilot, Claude. Acronyms are defined where they first appear and in the glossary (appendix).

AI assistants now answer buying questions, and three things are sold around the answer

Artificial intelligence (AI) assistants — ChatGPT, Gemini, Perplexity, Copilot, Claude — increasingly give the answer a shopper used to find by searching and clicking. The question this research asks: **does real demand exist, segment by segment, for brand visibility and product recommendation inside those answers, and what shows it works?**

Organic recommendation

The brand is named in the answer, unpaid. Sold as tooling and services to raise and measure mentions (also called GEO, generative engine optimisation).

Buyer: brand marketing, search owner

Paid placement

Ads shown inside or beside the AI answer, sold by the engine.

Buyer: media buyer

Agentic commerce

The assistant itself completes the purchase (checkout inside the assistant).


Buyer: e-commerce, payments

The three are kept separate throughout: their sizes, prices and evidence are never added together.

WHY NOW

Between 2026-02-09 and 2026-09-17, engines, investors, regulators and courts all acted

Selected events, in date order; spacing is not to scale



The record holds **54** dated events, 2023-04-25 to a 2027-01 compliance date. The open question for a director: **do brands want this, and does it pay?**

As of 2026-09-23 the evidence shows demand at the level of job postings and ad inventory, and none at the level of payback

1 · No sub-market has a measured size

Only floors from disclosed revenue, and forecasts that disagree widely.

2 · Many sellers, little disclosure

Vendors and engines rarely publish prices, revenue or method.

3 · Buyers spend in 8 of 27 segments

Mostly shown by job postings; no buyer discloses a price paid.

4 · No causal case; filed evidence runs negative

No published case isolates a sales effect.

Then: eleven ranked risks, and where the evidence stops.

Every figure on the slides that follow carries a source kind, a tier from 1 (strongest) to 7 (weakest), and a date

Tier: provenance, 1 strongest to 7 weakest

Tier	What it is
2	Filed: regulatory, court or securities filing
3	Platform primary: the engine or company's own statement
4	Panel, clickstream, telemetry; preprints
5	Vendor or agency study with sample size, dates, method
6	Marketing, "Not a source"

Source kind

vendor-reported · analyst-derived · filed · company-stated · measured-by-us (our own count or screen)

Case grades for "does it work"

Grade	Design
Gold	Holdout, geo-split or switchback: causal
Silver	Before/after with an unaffected control: correlational
Bronze	Visibility change only

Silver is counted two ways: **as graded**, and under the **literal rule** (capped at Bronze when any required item is missing).

† on a slide: the figure has another reading or date; both stand in the appendix, never averaged.

None of 18 published size figures measures a sub-market; 15 are forecasts

Count of published size figures



Linear scale from 0.

The source table grades this set of size figures tier 6: marketing-grade.


What can be stated instead: a **floor** from revenue that sellers disclose, and the **spread** between forecasts. Both follow on the next slide.

Disclosed revenue gives floors only; forecasts for each sub-market disagree by 1.92× to 26.3×

Sub-market	Floor from disclosed figures	Forecast range (not a size)
Organic	\$42.2M–\$48.2M + €2.2M ARR, 4 vendors, 2025-11 to 2026-05 +	2034: \$17.15B–\$32.92B
Paid	≥ \$1B run rate, one engine, 2026-08-31 +	2030: ~\$5B–\$100B+
Agentic	unknown — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon 2026-09-23	2030: \$190B–\$5T

ARR: annual recurring revenue. Run rate: current revenue annualised, not booked revenue. Currencies left as stated.


Forecast spread, high ÷ low, as the sources state it



Linear scale from 0x. Ratios, not sizes.

34 organic-visibility vendors are rostered; 14 publish a price and none publishes its prompt set

Rostered vendors, count out of 34



Linear scale from 0.

- Published self-serve prices: "\$20/mo" to "\$999/month".
- The 2 filers do not break this product out (FY2026; FY: fiscal year).
- Filed product line: Semrush, "AI products surpassed \$38 million in ARR", 2025-12.
- No vendor publishes the fixed set of prompts, with sample size n, behind its visibility scores.

Four engines sell ads inside answers, but none of 8 engines publishes a rate card

Selling ads in answers as of 2026-09-23: **ChatGPT, Google, Copilot, Amazon Rufus** †

Disclosure, of 8 engines	Count
Publish a rate card	0 of 8
Disclose ad revenue	1 of 8

Only engine-stated price: OpenAI's maximum-bid guidance, "\$3–\$5 USD per click" (USD: US dollars).

ChatGPT ads revenue: "\$1 billion in annualized revenue run rate", 2026-08-31 †

How many advertisers? Two sources disagree

"tens of thousands"

OpenAI, 2026-08-31
company-stated, tier 3

820 to 7,378

Ad-tracking panels, 2026-07
vendor-reported, tier 5

Shown side by side; neither is preferred. An earlier count and dated shares of ad presence are in the appendix.

company-stated, tier 3, 2026-08 and 2026-09; vendor-reported, tier 5, 2026-07

Agentic checkout fees are mostly undisclosed: 1 of 4 live engine programmes states one

1 of 4

live engine checkout programmes states its fee

Copilot Checkout "does not take a commission or affiliate fee" (2026-09) †


6 of 6

checkout-protocol fee clauses: unknown — checked 2026-09-22

What a director would need	Status
Revenue floor	unknown — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon 2026-09-23
Fee per transaction	stated by 1 of 4 programmes
Forecast	2030: \$190B–\$5T, analyst-derived, tiers 5–6 †

8 of 27 buyer segments show spend, under either reading of unchecked channels

27 segments by read, 2026-09-23 (count of segments)



Segment = sub-market × vertical (skincare, B2B SaaS, high-CPA regulated) × buyer size (SMB, mid-market, enterprise).


B2B SaaS: business-to-business software sold as a service. High-CPA regulated: credit cards, insurance, supplements; CPA is cost per acquisition. SMB: small and medium business.

Strict: a still-blocked channel leaves the segment unchecked (blank). **Loose:** a blocked-and-logged channel counts as checked. Attention: interest without spend.

† An earlier 2026-09-22 reading is in the appendix.

Spend sits mostly in organic tooling and large buyers, and shows mostly through job postings

Spend segments by sub-market, 2026-09-23



Count of segments reading spend; linear scale from 0.

On 2026-09-22 (earlier reading, 7 spend segments)

- 6 of the 7 rested on job postings — the weakest spend signal (company-stated, tier 3)
- 5 of the 7 were enterprise buyers
- 1 of 9 small and medium business (SMB) segments read spend

On 2026-09-23

- Willingness to pay: unknown in all 27 segments
- Local businesses, 9 segments read separately: spend 1, attention 3, none — checked 5

No published case proves a causal sales effect: zero Gold among ~980 and ~1,640 more screened cases

Gold (causal)
none in ~980 cases to 2026-09-22; none in ~1,640 more on 2026-09-23

Silver (before/after with control)
2026-09-22: 7 as graded, 1 literal; 2 of the 7 run negative or null
2026-09-23: +4 as graded, +1 literal; 0 new Silver in any vertical

Brand confirms the vendor's case
1 of 168 checked, cumulative to 2026-09-23 + (measured-by-us)

Paid ads results with a control
0 of 15 screened (measured-by-us, 2026-09-23)

Cases come from vendors (tier 5). The strongest grade found is correlational, and nearly no brand confirms its vendor's case.

SE Ranking's own ChatGPT ad campaigns: click-through rate (CTR) 1.30% on 97,000+ impressions, "very few sign-ups" (vendor-reported, tier 5, 2026-08-10).

The strongest filed evidence runs negative for publishers and comparison sites

"decreased 24%"

NerdWallet credit-cards revenue
filed, tier 2, 2026-02-25

"63% decline"

People Inc Google Search referrals
"over two years"
filed, tier 2, 2026-05-04

"83-93% drops"


Click-through rates for The Times and
a co-plaintiff publisher group,
recorded by Microsoft †
filed court brief, tier 2, 2026-09-17

Null result: generative engine optimisation (GEO) levers moved citation on "none" of ten engine families tested (analyst-derived preprint, tier 4, 2026-09-07).

Revenue, referrals and click-through are different metrics, shown separately and not compared. No causal link to AI answers is drawn here.

Brands that moved changed content and adopted vendor tooling; 2 of 99 changed paid media

Brands naming each change (of 99 naming any)



108 brands moved; 99 name a change. Linear scale from 0.

Evidence class of the 108 cases, two readings

Class	Experimental reading	Observational reading
Experimental	6	1
Observational	61	66
Action only	37	37
Outside	4	4

The readings differ on one design: a before/after with a control counts as experimental in the first, observational in the second.

Referral collapse and regulation lead eleven ranked risks; no source states likelihood or impact

Rank	Risk	Strongest evidence
1	Referral collapse: the answer replaces the click	Microsoft, NerdWallet filings (filed, tier 2)
2	Regulation: disclosure duties, fair-ranking orders	AI Act; search-engine designation; UK CMA order 2026-06-17 (filed, tier 2)
3	Vendor consolidation	Adobe–Semrush; Yext–GoShine 2026-09-01 (filed, tier 2)
4	Platform capture: engines own ads, checkout, measurement	Google: "you can't opt out" (company-stated, tier 3)
5	No payback evidence	0 Gold; 0 of 15 ad results controlled
6	Measurement opacity	ChatGPT ad presence 0.8%–26% + (vendor-reported, tier 5)
7	Substitutes suffice	Google: "no additional requirements... nor other special optimizations necessary" (tier 3)
8–11	Engine policy reversal; agentic fees; manipulation; unsizeable market	

Likelihood, impact, mitigation: no source states one. UK CMA: United Kingdom Competition and Markets Authority.

The evidence stops at price paid, conversion, and what engines recommend unprompted

Question	Status
A price any buyer paid	unknown in all 27 segments, 2026-09-23
Does AI referral convert above organic search?	unknown — checked Adobe, Salesforce, Shopify 2026-09-22
What engines recommend under neutral conditions	unknown — not in repo; our own sampling stopped partway, held 2026-09-22 22:40

Channels still closed at the 2026-09-23 re-probe

- Blocked: indeed.com, emarketer.com, perplexity.ai/hub
- Payment required: Crunchbase, Gartner, Forrester
- Reddit readable only through a public archive

Any of these could move a segment from "none — checked" or blank to spend.

Demand is real enough to hire and to buy inventory; no evidence yet shows it pays back

Your question	What the evidence supports	Where it stops
Why now?	Ads in answers since 2026-02-09; regulation and court filings through 2026-09-17	—
How big?	Floors only; forecasts disagree by up to 26.3×	No measured size in any sub-market
Who is there?	34 organic vendors; four engines selling ads	Prices and fees mostly undisclosed
Who buys?	8 of 27 segments show spend, mostly job postings	Price paid unknown in all 27
Does it work?	0 Gold; filed evidence negative	No causal case either way

No decision is requested and no lean is attached; the verdict is yours. These reading choices move figures, each shown both ways in the appendix: Silver as graded or literal; strict or loose; experimental or observational.

Glossary

- **AI** — artificial intelligence. **AI Act Art. 50** — EU transparency obligations for AI systems.
- **EU / UK / US** — European Union / United Kingdom / United States. **USD** — US dollars. **FY** — fiscal year.
- **ARR** — annual recurring revenue. **Run rate** — current revenue annualised, not booked.
- **B2B SaaS** — business-to-business software as a service. **High-CPA regulated** — credit cards, insurance, supplements; CPA: cost per acquisition. **SMB** — small and medium business.
- **CMA** — UK Competition and Markets Authority. **Very Large Online Search Engine** — EU Digital Services Act designation, carrying an ad-repository duty from 2027-01 for ChatGPT.
- **CTR** — click-through rate. **GEO** — generative engine optimisation: content tactics aimed at being cited by AI answers.
- **Segment** — sub-market × vertical × buyer size, read spend, attention, none — checked, or blank (not checked).
- **Strict / loose** — strict: a still-blocked channel leaves the segment unchecked; loose: a blocked-and-logged channel counts as checked.
- **Gold / Silver / Bronze** — case grades: causal / before-after with control / visibility change only. **As graded / literal rule** — as assigned / capped at Bronze when a required item is missing.
- **Experimental / observational reading** — whether a before/after with a control counts as experimental or observational.

Market and seller figures with a second reading

Figure	Reading A	Reading B
Organic floor	"\$5.3M–\$35.2M annualised", price × count, vendor-reported, tier 5, 2026-09	"\$42.2M–\$48.2M + €2.2M ARR", disclosed ARR, tiers 2/5, 2025-11 to 2026-05
Agentic forecast spread	"roughly 35×", 2029–30	26.3×, 2030 only
Paid engines live	"5 of 8", undated	"Four engines", 2026-09-23
ChatGPT advertisers	"over 600", tier 5, 2026-03-26	"tens of thousands", tier 3, 2026-08-31; panels 820 to 7,378
ChatGPT ad presence	0.8% of 500+ prompts, launch weeks	4.47% US queries 2026-03 to 2026-05; "0.00 percent" UK 2026-06; 24.7% 2026-07; 25.94% of 50,006 US prompts, published 2026-08-10
ChatGPT ads revenue	"\$1 billion in annualized revenue run rate", tier 3, 2026-08-31	"roughly \$83 million a month", tier 5, 2026-08-31
Microsoft CTR drops	"83-93%", Times and co-plaintiff, brief p.10, 2026-09-17	"87% to 93%" Times, "83% to 91%" co-plaintiff, brief p.76
Copilot Checkout fee	0%, "does not take a commission"	unknown in a protocols-table pull

Demand, proof and funding figures with a second reading

Figure	Reading A	Reading B
27-segment tally	spend 7 / attention 1 / none 11 / blank 8, 2026-09-22	2026-09-23: strict 8 / 1 / 10 / 8, loose 8 / 1 / 18 / 0; after re-probe strict 8 / 1 / 17 / 1, loose 8 / 1 / 18 / 0
Silver count	7 as graded, 1 literal, 2026-09-22	+4 as graded, +1 literal, 2026-09-23
Cases screened, no Gold	~980, 2026-09-22	~1,640 more, 2026-09-23
Brand confirmation	0 of 59, 2026-09-22	1 of 109; cumulative 1 of 168, 2026-09-23
Skincare × agentic × enterprise	spend, undated	attention, undated
Scrunch deal value	"not disclosed", tier 3, 2026-06-03	"\$225 Million", tier 5, 2026-06-03
Peec AI ARR	"more than \$4 million", 2025-11-17	"10m", 2026-05-23
geoSurge seed	"\$12M seed", 2026-07-07	"€10M", same round, 2026-07-07
Sitefire / Jerry case grade	Bronze in one case count, 2026-09-22	Silver in another, 2026-09-22
Share of key claims on tier-3 sources	17 of 25 = 68.0%, 2026-09-22	2026-09-23 recounts: 12 of 27 = 44.4%; 11 of 25 = 44.0%; 12 of 19 = 63.2%; 12 of 16 = 75.0%; 12 of 24 = 50.0%