

Brand visibility inside AI assistants: is there demand, and does it pay?

What the published evidence shows about brands being recommended, advertised and sold inside artificial intelligence (AI) assistants: ChatGPT, Google, Microsoft Copilot, Perplexity and Claude.

Research only: evidence, not a verdict

Evidence as of 2026-09

AI assistants now answer buying questions; the open question is whether brands will pay to appear, and whether it works

What is happening

AI assistants (ChatGPT, Google, Microsoft Copilot, Perplexity, Claude) answer buying questions in prose. A brand appears in those answers through three markets:

- **Organic recommendation:** unpaid mention, bought as software and agency work.
- **Paid placement:** ads inside the answer.
- **Agentic commerce:** the assistant completes the checkout.

What changed

In 2026 the engines switched on ads and checkout inside their answers, and vendors raised money on the promise. Sellers being live is not proof that buyers pay, or that paying works.

The question this deck answers

Does real demand exist, per buyer segment, and does anything show it pays?

Sellers are live and funded, and large companies are hiring; no published evidence shows it pays

What the evidence supports: supply exists and some buyers are spending. Where it stops: no market has a measured size, willingness to pay is unknown in all 27 buyer segments, and no published case proves payback.

WHY NOW Engines switched supply on in 2026 ChatGPT ads: "\$1 billion" annualized run rate (company-stated, 2026-08-31).	HOW BIG Floors and forecasts, no measured size 18 published size figures: 15 forecasts, 0 measured sizes (analyst-derived, 2026-09-22).	WHO SELLS Engines own the rails; vendors funded, disclosure thin One vendor raised "\$180M Series D at \$1.8B valuation" (vendor-reported, 2026-09-15).	WHO BUYS Enterprises, mostly by hiring; price paid unknown 8 of 27 buyer segments show spend, strict reading (analyst-derived, 2026-09-23).	DOES IT WORK No controlled result; best-sourced results negative No published case of ~980 screened (2026-09-22), nor ~1,640 more (2026-09-23), uses a withheld control group.
---	---	---	---	--

HOW TO READ THE EVIDENCE

Every figure carries who said it, when, and how strong the source is

Source kind	Meaning
filed	Regulator, court or statutory company filing
company-stated	A company's own page or release about itself
vendor-reported	A seller of the service reporting results
analyst-derived	Analysis over published data, ours or a third party's
measured-by-us	Our own count of published material

Tier ranks provenance, 1 strongest, 7 weakest: 1 measured by us, 2 filed, 3 platform's own page, 4 panel data, 5 vendor study, 6 marketing, 7 do not use.

Evidence grade	Meaning
Gold	Causal: holdout, geo-split or switchback design
Silver	Before/after with an unaffected control
Bronze	Correlational, or visibility change only

Holdout: a withheld control group. Geo-split: treated versus untreated regions. Switchback: treatment alternated over time.

Segment read	Meaning
spend	Money moved
attention	Interest without spend
none — checked	Every signal checked, none found
blank	Not yet checked

In 2026 the engines switched on checkout and ads, and regulators followed

2026-01-08	○	Copilot Checkout live, United States (US); "does not take a commission" company-stated
2026-02-04	○	"Claude will remain ad-free" company-stated
2026-02-09	○	ChatGPT Ads live, US pilot company-stated
2026-06-17	○	United Kingdom (UK) fair-ranking requirement covers "search generative AI features" filed
2026-08-31	○	ChatGPT Ads "\$1 billion... run rate"; >1B weekly users company-stated
2026-08-31	○	European Union (EU) designates ChatGPT a Very Large Online Search Engine, 159.1M EU recipients; ad-repository duty 2027-01 filed
2026-09-17	○	Court exhibit, Microsoft: "83-93% drops in click-through rates" filed

54

dated events, 2023-04-25 to 2027-01

26 at tier 2 (filed)

24 at tier 3 (platform's own page)

Selected events shown, in date order. Measured-by-us, 2026-09-23.

No market has a measured size: of 18 published size figures, 15 are forecasts and 0 are measured

Sub-market	Disclosed floor (not a size)	Forecast spread (analyst-derived, not a size)
Organic	\$42.2M–\$48.2M + €2.2M annual recurring revenue (ARR), 4 vendors, 2025-11 to 2026-05 (filed; company-stated). Another build exists (appendix)	2034: \$17.15B–\$32.92B, 1.92×, 4 forecasts, tier 6
Paid	≥ \$1B run rate, one engine, 2026-08-31 (company-stated); 0.45% of Google Search & other 2025, bases differ	2030: ~\$5B–\$100B+, ~20×, scopes differ
Agentic	unknown — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon 2026-09-23	2030: \$190B–\$5T, 26.3×, 7 forecasts; another reading exists

2.5B

monthly users of Google AI Overviews (company-stated, 2026-05-19)

1.65%

AI tools' share of US desktop visits, first quarter (Q1) 2026 (analyst-derived)


76%

of "more than 200" ad buyers name AI answers their top focus: attention, not spend (analyst-derived, 2026-09-10)

HOW BIG

Forecasts do not agree: high and low differ by 1.92× in organic and by 26.3× in agentic

Spread: highest forecast ÷ lowest forecast, times (×). Linear scale, 0 to 40×. Ratios as stated in the source.



The two agentic readings sit side by side; neither is picked. Forecast dollar ranges are for different years and scopes and are not compared.

The engines own the rails: four sell ads, none publishes a price list, one discloses a checkout fee

Ads inside answers

- **Four** engines sell them: ChatGPT, Google, Copilot, Amazon Rufus (company-stated, 2026-09-23). Another count exists: 5 of 8 with a live product.
- **0 of 8** publish a rate card (a published unit price list).
- Only stated price: OpenAI's "\$3–\$5 USD per click" bid guidance (company-stated).
- Claude refuses; Perplexity is winding ads down "by the end of 2026" (analyst-derived relay).

Checkout inside answers

- **4** live checkout programs (company-stated, 2026-09).
- **1 of 4** discloses a fee: Copilot, 0% (company-stated, 2026-09).
- **6 of 6** protocol fee clauses: unknown — checked 2026-09-22.

Rufus: Amazon's shopping assistant. Protocol: the shared technical rules an assistant uses to pay a merchant.

34 organic-recommendation vendors are rostered, and most disclose little about customers or revenue

34


vendors rostered (measured-by-us, 2026-09)

13 pure-play

16 incumbent bundles

5 agencies

Semrush: "AI products surpassed \$38 million in ARR" (filed, 2026-03-02)



Filed revenue: neither of the 2 breaks out the AI product. Prompt set: the questions a vendor tests engines with.


Money is moving in: incumbents are buying vendors, and one vendor is valued at \$1.8B

- 2026-04-28 ○ Adobe–Semrush acquisition filed
- 2026-06-03 ○ Sitecore–Scrunch acquisition. Value: **not disclosed** by Sitecore (company-stated) | **\$225M** per Bloomberg (press relay), unconfirmed
- 2026-09-01 ○ Yext–GoShine acquisition filed
- 2026-09-15 ○ One vendor raises "\$180M Series D at \$1.8B valuation" vendor-reported

The two Scrunch values sit side by side: the buyer states none, a press report states one. Neither is picked.

Spend shows in 8 of 27 buyer segments under either reading; 17 or 18 show none after checking

Buyer segments by read, of 27 (analyst-derived, 2026-09-23)
Solid = strict reading; outlined = loose. Bar length proportional to count.



What a segment is

Sub-market × vertical × buyer size.
Verticals: skincare and beauty; business-to-business software (B2B SaaS); high cost-per-acquisition (CPA) regulated.

Spend by sub-market

Organic 6, paid 1, agentic 1; that agentic cell is also read as attention (analyst-derived, 2026-09).

Strict: a blocked channel does not count as checked. An earlier reading (2026-09-22) exists: appendix.

Spend comes from enterprises, mostly seen as job postings; the price buyers pay is unknown

Buyer size and signal (2026-09-22)

- **5 of the 7** spend cells were enterprise.
- Exactly **1 of 9** small and medium business (SMB) cells read spend.
- **Six of the seven** rested on job postings, the weakest spend signal (company-stated).

A fourth vertical: local and multi-location

9 cells: spend 1, attention 3, none — checked 5 (analyst-derived, 2026-09-23).

The one buyer stating paid spend

A Spanish supplement brand: "**€37.61**" on ChatGPT Ads, "**0 conversions**" (company-stated, tier 5, 2026-09-23).

Willingness to pay

unknown — not in repo, in all 27 segments.

No published case uses a causal design; Silver results shrink under a stricter grading rule

0	Gold: holdout, geo-split or switchback Vendor-reported cases, 2026-09-23
7 as graded 1 under rule 1	Silver: before/after with a control Vendor-reported cases, 2026-09-22. A re-pass adds +4 as graded and +1 under rule 1 (2026-09-23).
count not shown	Bronze: correlational, or visibility change only

What was screened
~980 published cases (vendor-reported, 2026-09-22), and ~1,640 more (2026-09-23). None uses a holdout, geo-split or switchback.

Grading rule 1
A case missing any of brand, engine, absolute date window, baseline, intervention, sample size, or who measured is Bronze at best.

The best-sourced results run negative or null, and brands almost never confirm vendor claims

2 of 7 Silvers run negative or null

- NerdWallet: cards revenue "decreased 24%" (filed, 2026-02-25).
- Generative engine optimisation (GEO) levers move citation on "none" of ten engine families (preprint, 2026-09-07).

Brand confirmation

Brands confirming a vendor case on their own domain: cumulative **1 of 168**, Chime (measured-by-us, 2026-09-23). Another reading exists: appendix.


Paid ads

- **15** results screened, **0** cleared.
- **10** observational advertiser reports, **0** controlled (vendor-reported, 2026-09).
- One campaign: click-through rate (CTR) 1.30% on 97,000+ impressions, "very few sign-ups" (vendor-reported, 2026-08-10).

108 early movers changed content and tooling; almost none changed paid media, and almost none ran an experiment

What they changed, of 108 movers

Measured-by-us over self-reports, 2026-09. Bars to scale.




Both paid-media moves: Google's Direct Offers pilot.

R1 counts before/after-with-control cases as experimental; R2 counts vendor-run ones as observational. Both readings shown; neither picked.

Evidence class, two readings (2026-09-23)

Solid = R2; outlined = R1. Bar length proportional to count.



Eleven risks, ranked by strength of evidence; referral collapse carries the strongest

Rank	Risk	Strongest evidence
1	Referral collapse	Court exhibit; NerdWallet; People Inc "22% decline in Core Sessions" (filed, 2026-08-03)
2	Regulation	EU AI Act Art. 50 in force 2026-08-02; search-engine designation; UK order (filed)
3	Vendor consolidation	Acquisitions on slide 10 (filed)
4	Platform capture	OpenAI Pixel, Conversions API; Copilot fee (company-stated)
5	No payback evidence	0 Gold; 0 of 15 ad results controlled (vendor-reported, 2026-09-23)
6	Measurement opacity	ChatGPT ad presence 0.8%–26% (vendor-reported, pulled 2026-09-23)
7	"Do nothing" suffices	Google: "no additional requirements"
8	Engine policy reversal	Claude; Perplexity; Gemini "not rushing anything here"
9	Agentic rails	Fees undisclosed; merchants auto-enrolled
10	Manipulation and countermeasures	Google spam policy, 2026-08-28 (company-stated)
11	Unsizeable market	Forecasts only

Ranked by tier of strongest evidence, then breadth; not by severity. Likelihood, impact, mitigation, owner: no source states one.

The evidence cannot show sales lift, conversion versus search, or what engines actually answer

Not shown by any source

- Sales lift from any intervention.
- Whether AI referral out-converts organic search: unknown — checked Adobe, Salesforce, Shopify 2026-09-22.
- What engines return to buying prompts: our own sampling was halted part-way, 2026-09-22.

Channels we could not read

- Blocked 2026-09-23: reddit.com, indeed.com, emarketer.com.
- Paywalled: Crunchbase, Gartner, Forrester ("\$1495" per report, company-stated).

Readings the evidence does not settle: each option, with its evidence

1. Which sub-market a hypothetical product would target, or more than one

Organic: 6 spend cells; 34 vendors; floor \$42.2M–\$48.2M + €2.2M ARR.

Paid: 1 spend cell; ≥ \$1B run rate at one engine; 0 of 15 ad results controlled.

Agentic: 1 spend cell, also read as attention; 1 of 4 checkouts discloses a fee.
2. Which Silver count governs your reading of proof

As graded: 7 (2026-09-22), +4 on re-pass. **Under grading rule 1:** 1, +1 on re-pass. Gold is 0 under both.
3. Which demand reading you accept

Strict: spend 8, attention 1, none — checked 17, blank 1. **Loose:** spend 8, attention 1, none — checked 18, blank 0. Spend is 8 of 27 under both.

Both readings stand; none averaged or picked

Figure	Reading A (date)	Reading B (date)
Demand tally, 27 segments	spend 7 / attention 1 / none — checked 11 / blank 8 (2026-09-22)	strict 8/1/17/1, loose 8/1/18/0 (2026-09-23)
Silver count	7 as graded, 1 under rule 1 (2026-09-22)	+4 as graded, +1 under rule 1 (2026-09-23)
Brand confirmation	0 of 59 (2026-09-22)	1 of 109; cumulative 1 of 168 (2026-09-23)
108-mover evidence class	R1: 6 / 61 / 37 / 4 (2026-09-23)	R2: 1 / 66 / 37 / 4 (2026-09-23)
Agentic forecast spread	26.3×, 2030, 7 forecasts (date not stated)	"roughly 35×", 2029–30 forecasts (date not stated)
Engines selling ads	Four, own pages (2026-09-23)	5 of 8 live product (date not stated)
Organic floor	\$42.2M–\$48.2M + €2.2M ARR, disclosed ARR (2025-11 to 2026-05)	\$5.3M–\$35.2M, price × count (2026-09)
Skincare × agentic × enterprise	spend (date not stated)	attention, "below the tier-5 spend floor" (date not stated)
Scrunch deal value	not disclosed, Sitecore (2026-06-03)	\$225M Bloomberg, unconfirmed (date not stated)

Terms and acronyms

Term	Meaning
AEO / GEO / SEO	Answer engine / generative engine / search engine optimisation. GEO never means geographic here
AI; API	Artificial intelligence; application programming interface, a machine-to-machine data feed
AI Overviews; AI Mode	Google's AI answer block on a results page; Google's conversational search tab
ARR	Annual recurring revenue
B2B SaaS	Business-to-business software sold as a subscription
CPA; high-CPA regulated	Cost per acquisition; the vertical where acquiring a customer is costly and regulated
CTR	Click-through rate: clicks divided by impressions
Direct Offers	Google's AI Mode ad pilot, unit labelled "Sponsored deal"
EU AI Act Art. 50	EU transparency duties for AI systems; names no assistant
Grading rule 1	A case missing any of brand, engine, absolute date window, baseline, intervention, sample size, who measured is Bronze at best

Terms and acronyms

Term	Meaning
Pixel; Conversions API	OpenAI's advertiser tracking tag and server-side conversion feed
Q1	First calendar quarter
R1 / R2	Two readings of before-after-with-control cases: R1 counts them experimental, R2 counts vendor-run ones observational
Rate card	A published unit price list
Rufus	Amazon's shopping assistant
SMB	Small and medium business
Strict / loose	Whether a blocked, logged channel counts as checked when a segment reads "none"
Tier	Source provenance: 1 measured by us, 2 filed, 3 platform's own page, 4 panel data, 5 vendor study, 6 marketing, 7 do not use
US; UK; EU; USD	United States; United Kingdom; European Union; US dollars
VLOSE	Very Large Online Search Engine: an EU Digital Services Act designation that brings ad-repository duties