

BRAND VISIBILITY AND PRODUCT RECOMMENDATION INSIDE AI ASSISTANTS

Brands have started paying to be found inside AI assistants; the published record stops short of showing it pays

Research only: evidence, not a verdict

Evidence as of 2026-09

Assistants in scope: ChatGPT, Gemini, Perplexity, Copilot, Claude. AI: artificial intelligence. The programme produces evidence; the verdict belongs to the reader.

Engines now sell space inside answers, few buyers spend, and no published case shows it pays

What the evidence supports: spend has begun. Where it stops: market size and payback. The verdict is yours.

Supply is live

Four

engines sell ads inside AI answers
(*company-stated, 2026-09-23*).

ChatGPT ads reached a "\$1 billion"
run rate (*company-stated, 2026-08-31*).

Demand is thin

8 of 27

buyer segments show a spend signal
(*measured-by-us, 2026-09-23*), most
resting on job postings, not category
budgets.

No sub-market has a measured size.

Proof is missing

None

of the published cases has a causal
design (holdout or split test).

The strongest filed evidence runs
negative: "83-93% drops in click-
through rates" (*filed, 2026-09-17*).

The field splits into three markets, and visibility, traffic and sales are never counted as one


Sub-market	What is sold	Who buys
Organic recommendation also called generative engine optimisation (GEO), answer engine optimisation (AEO), "AI visibility"	Help for a brand to appear, unpaid, in an AI answer: tools, audits, agencies	Brand marketing; the search engine optimisation (SEO) owner
Paid placement AI ads	Ad units shown inside AI answers, sold by the engine	Media buyer
Agentic commerce	Checkout run by the assistant on the shopper's behalf	E-commerce and payments teams

Visibility The brand is named in an answer.	Traffic The answer sends visits to the brand's site.	Sales Those visits buy. A visibility gain is never read as sales proof.
---	--	---

Each number here is labelled by who produced it; lower tier is stronger, and only Gold means causal

Source kind: who produced the number		Tier	Provenance
Filed	Regulatory or court filing	1	Measured by this programme
Company-stated	A company about itself, outside a filing	2	Filed: annual report, court exhibit, funding filing
Vendor-reported	A seller about its own product or market	3	A platform's own docs, changelog or pricing page
Analyst-derived	A third party computed it	4	Panel, clickstream, telemetry
Measured-by-us	This programme's own pull or count	5	Vendor or agency study with sample size and method
Gold: causal design (holdout, geo-split, switchback). Silver: before-and-after with an unaffected control; correlational. No Gold case exists.		6	Vendor blog, no sample size: marketing, "not a source"
		7	Listicle or roundup: not used

Engines began selling ads inside answers in 2026-02; regulators and courts followed by 2026-09



Filled dot: engines, vendors, buyers of vendors. Hollow dot: regulators and courts. Events in date order; spacing not to scale.

Four engines sell ads inside answers, yet none of eight publishes a rate card

Engine	Status and stated terms
ChatGPT (OpenAI)	Selling ads; the only stated price is "\$3–\$5 USD per click" bid guidance (<i>company-stated, 2026-09-22</i>); own Ads Manager and Pixel
Google AI Overviews and AI Mode	Selling ads
Microsoft Copilot	Selling ads; its checkout "does not take a commission" (<i>company-stated, 2026-09</i>)
Amazon Rufus	Selling ads
Claude (Anthropic)	Will "remain ad-free" (<i>company-stated, 2026-02-04</i>)
Perplexity	"Winding down" ads "by the end of 2026" (<i>analyst-derived relay, 2026-02-18</i>)

A second count on another engine base, 5 of 8 engines with a live product, also stands (appendix).

0 of 8

engines publish a rate card (*company-stated, 2026-09-22*)

1 of 4


live assistant checkout programmes discloses a fee (*company-stated, 2026-09*)

34 vendors sell organic visibility; most disclose neither price nor revenue nor method


Who the 34 are (analyst-derived, 2026-09)




What they disclose, of 34



"\$20/mo" to "\$999/month" (vendor-reported, 2026-09)



neither breaks out AI visibility (filed, FY2026)



(vendor-reported, checked 2026-09-22)

Bar length: count of vendors; full grey bar = 34.

Semrush: "AI products surpassed \$38 million in ARR" (annual recurring revenue) *(filed, 2025-12)*

Profound: "\$180M Series D at \$1.8B valuation" *(vendor-reported, 2026-09-15)*

Sitecore bought Scrunch, 2026-06-03. Value "not disclosed" *(company-stated)* beside a "\$225 Million" press report *(analyst-derived)*

No one has measured the size of any of the three markets; only disclosed revenue floors exist


Sub-market	Disclosed floor
Organic	\$42.2M–\$48.2M + €2.2M ARR, 4 vendors, 2025-11 to 2026-05 <i>(filed; company-stated)</i> A price × customer build gives \$5.3M–\$35.2M <i>(vendor-reported inputs, 2026-09)</i>
Paid	≥ \$1B run rate, ChatGPT ads alone, 2026-08-31 <i>(company-stated)</i>
Agentic	unknown — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon, 2026-09-23

M / B: million / billion. Floors are what named sellers disclose, not market totals.

O measured market sizes among **18** published figures, of which **15** are forecasts *(analyst-derived, 2026-09-23)*


Published forecasts for one market disagree by up to 26.3×, roughly 35× on a second reading; none is a size

Spread: highest forecast ÷ lowest, as computed in the source files (unit: ×)




8 of 27 buyer segments show spend; 17 show nothing on the strict reading

Strict reading, 2026-09-23



Attention 1 (hatched) · Blank 1 (dashed)

Loose reading, 2026-09-23



Attention 1 (hatched) · Blank 0


Bar length: count of cells; full bar = 27 cells

A segment is one cell: sub-market × industry (skincare; business software; high-cost regulated) × buyer size (small, mid-market, enterprise).

Spend money moved; **attention** interest only; **none** every signal checked, nothing found; **blank** not yet checked. **Strict / loose:** whether a still-blocked channel counts as checked.

Spend by market: organic 6, paid 1, agentic 1; a second compiler reads that agentic cell as attention.

The spend signals rest mostly on enterprise job postings, and no buyer's price is known



Wider checks, measured-by-us, 2026-09-23	Result
Local and multi-location chains, nine cells	1 spend, 3 attention, 5 none
UK, France, Spain, Italy, Netherlands	2 spend of 45 cells
One buyer's own report, Spanish supplement brand	"€37.61" spend, "0 conversions" (<i>company-stated, 2026-09-23</i>)

Screening the published cases found no causal result; the Silver count depends on the rule applied

Cases screened

No Gold

~980 cases on 2026-09-22 and
~1,640 more on 2026-09-23; no Gold
(causal) case in either (*measured-by-us*)

Silver, 2026-09-22

7 or 1

7 by census grade; 1 under the strict
item rule. 2 of 7 are negative or null.
A 2026-09-23 re-run added +4 and
+1 (*measured-by-us*)


Brand confirms vendor

1 of 168

brands confirm a vendor's case on
their own site, cumulative, 2026-09-
23. An earlier 0 of 59 also stands
(*measured-by-us*)

Census grade: the grade the screen assigned. **Strict item rule:** capped below Silver when any of seven items (brand, engine, date window, baseline, intervention, sample size, who measured) is missing.

Of 108 cases where brands moved into AI answers, 1 shows an experimental result on reading R2, 6 on R1



Reading R1 (light bars) counts a before-and-after with a control as experimental.

Reading R2 (dark bars) counts it as observational.

Both readings stand; the source does not pick one.

The strongest published evidence runs negative: AI answers cut clicks and revenue for publishers

Who	What the source states	Date	Source kind
Microsoft, court exhibit	"83-93% drops in click-through rates"	2026-09-17	filed
NerdWallet	Credit cards revenue "decreased 24%"	2026-02-25	filed
IAC / People Inc (US publisher group)	"63% decline in Google Search referrals over two years"	2026-05-04	filed
TW3 / Citead (replication study)	Published page-edit levers move citation on "none" of ten engine families	2026-09-07	analyst-derived
HubSpot	Blog "5 million" fewer visits, thirty days	2026-09-22	company-stated

Paid ads: 0 of 15 advertiser results were controlled
(measured-by-us, 2026-09-23)

What exists instead: 10 observational advertiser reports
(company-stated, 2026-09-23)

The best-evidenced risk is referral collapse: the AI answer replaces the click, on filed court evidence

Rank	Risk	Strongest evidence	Source kind, tier
1	Referral collapse	"83-93% drops in click-through rates"	filed, 2
2	Regulation: disclosure, exclusions, fair ranking	EU AI Act Art. 50; VLOSE designation; CMA order	filed, 2
3	Incumbents absorb the tools	Semrush to Adobe; Yext buys GoShine 2026-09-01	filed, 2
4	Engines sell ads, checkout, measurement themselves	OpenAI Ads Manager and Pixel; Copilot takes no commission	company-stated, 3
5	No payback evidence	No Gold; 0 of 15 ad results controlled; HubSpot "5 million" fewer visits	company-stated, 3

Ranks 6–11: measurement opacity; "do nothing" suffices; engine policy reversal; undisclosed agentic fees; manipulation; unsizeable market.

Ranked by the tier of the strongest evidence, not by severity. Likelihood, impact, mitigation, owner: no source states one.

Spend has begun, but size and payback stay unproven; three open readings decide how strong the case looks

Which reading governs

- Buyer segments: strict or loose count
- Silver cases: 7 or 1
- Experimental movers: 1 or 6 of 108

What no source settles

- Size of any market: unknown — checked research-house pages 2026-09-22
- Whether AI referrals convert better than search: unknown, checked Adobe, Salesforce, Shopify 2026-09-22
- Price any buyer paid

Where the record is thin

- Own sampling of assistant answers stopped 2026-09-22; skipped 2026-09-23
- Walled 2026-09-23: reddit.com, indeed.com, emarketer.com, ted.europa.eu
- Crunchbase, Gartner, Forrester need payment

The decision is the reader's. This deck supplies the evidence and its limits; it does not recommend proceeding or stopping.

Appendix A1: glossary

AEO / GEO: answer / generative engine optimisation; names for the organic market. GEO never means geographic here.

AI: artificial intelligence. **AI Overviews / AI Mode:** Google's answer block on a results page / its conversational search tab.

Ads Manager, Pixel: OpenAI's self-serve ad-buying tool and its conversion-tracking tag.

ARR: annual recurring revenue. **Run rate:** current revenue pace, stated yearly.

CMA: UK Competition and Markets Authority.

DSA: EU Digital Services Act. **VLOSE:** Very Large Online Search Engine, a DSA designation with an ad-repository duty.

EU / UK / US: European Union / United Kingdom / United States.

EU AI Act Art. 50: transparency article: notice of AI interaction, marking of synthetic content.

FY: fiscal year. **M / B / T:** million / billion / trillion. **USD:** US dollars.

Gold / Silver: evidence grades. Gold: causal design (holdout, geo-split, switchback). Silver: before-and-after with an unaffected control.

Census grade / strict item rule: the grade the screen assigned / a cap below Silver when one of seven evidence items is missing.

IAC / People Inc: a US publisher group. **TW3 / Citead:** authors of a published replication study.

Reading R1 / R2: a before-and-after with control counted as experimental / as observational.

Rufus: Amazon's shopping assistant. **SEO:** classical search engine optimisation. **Series D:** a late-stage venture round.

Spend / attention / none / blank; strict / loose: segment reads and reading rules, defined on slide 10.

Tier; source kinds: defined on slide 4.

Appendix A2: figures carried two ways, both standing (1 of 2)

Figure	Reading A	Reading B
27-segment tally	Seven spend, one attention, eleven none — checked, eight blank; 2026-09-22	Strict 8 / 1 / 17 / 1; loose 8 / 1 / 18 / 0; 2026-09-23
27-segment tally, re-run	Strict 8 / 1 / 10 / 8; 2026-09-23	Loose 8 / 1 / 18 / 0; 2026-09-23
Silver count	7 census grade, 1 strict rule; 2026-09-22	+4 census grade, +1 strict rule; 2026-09-23
Cases screened	~980; 2026-09-22	~1,640 more; 2026-09-23
Brand-side confirmation	0 of 59; 2026-09-22	1 of 109, cumulative 1 of 168; 2026-09-23
108 mover cases experimental / observational / action named / outside	R1: 6 / 61 / 37 / 4; 2026-09-23	R2: 1 / 66 / 37 / 4; 2026-09-23
52 graded cases experimental / observational / action named	R1: 10 / 29 / 13; 2026-09-23	R2: 5 / 34 / 13; 2026-09-23
Organic floor	\$5.3M–\$35.2M, price × customer count	\$42.2M–\$48.2M + €2.2M ARR, disclosed

Appendix A3: figures carried two ways, both standing (2 of 2)

Figure	Reading A	Reading B
Agentic forecast spread	"roughly 35×", 2029–2030 horizon	26.3×, 2030 only
Engines selling AI ads	Four; 2026-09-23 (company-stated)	5 of 8 engines carry a live product; date not stated
Scrunch deal value	"not disclosed" (company-stated)	"\$225 Million" (analyst-derived)
Skincare, agentic, enterprise	Spend	Attention
Profound round, source kind	Vendor-reported	Company-stated
ChatGPT ad presenece	4.47% of US queries, 2026-03 to 2026-05; 0.00 percent across 169,560 UK scrapes, 2026-06 (vendor-reported)	"26% of ChatGPT responses", no window stated (vendor-reported)
ChatGPT advertisers	"tens of thousands" (company-stated)	820 to 7,378, panels (vendor-reported)