

Brand visibility inside AI assistants: what the evidence shows

Do brands pay to be named when people ask ChatGPT, Gemini, Perplexity, Copilot or Claude what to buy, and does it work?

Research only: evidence, not a verdict

Evidence as of 2026-09

Brands are moving and engines are selling. No published document shows it pays.

Brands are hiring for AI visibility and engines sell ads inside answers, but no published document shows it pays

Demand is real but thin

Spend signals appear in **8 of 27** buyer segments under both readings of the evidence (2026-09-23).

Willingness to pay: **unknown in all 27.**

Proof stops before payback

We screened **~980** published cases on 2026-09-22 and **~1,640** more on 2026-09-23.

None compares results against a withheld group, region or period.

Size is forecast, not measured

Disclosed vendor revenue sets floors.

Every dollar market size on file is a **forecast**: 15 of 18 published figures, 0 measured sizes.

Where the evidence stops. It shows movement: hires, tools, ad programmes. It does not show sales lift, prices paid, or a measured market size.

What is yours to decide. Four readings of the same evidence change how strong it looks (last slide).

The market splits into three sub-markets, each bought by a different team

Organic recommendation

Brand named in AI answers, unpaid.
Called GEO (generative engine optimisation) or AEO (answer engine optimisation).
Buyer: marketing, or the SEO (search engine optimisation) owner.

Paid placement

Ads shown inside AI answers.

Buyer: media buyers.

Agentic commerce

The assistant runs the checkout for the user.

Buyer: e-commerce and payments teams.

Three metrics, never mixed: visibility (is the brand named), traffic (do users click through), sales (do they buy).

Verticals studied: skincare and beauty; B2B SaaS (business software sold by subscription); high-CPA regulated (cards, insurance, supplements; CPA: cost per acquisition); local multi-location businesses, added 2026-09-23.

Most figures here are self-reported, so each one carries a source kind and an evidence tier

Tier	What the evidence is (1 strongest, 7 weakest)
1	Measured by us
2	Filed (regulatory or court filing)
3	The platform's own page
4	Panel or telemetry data
5	Vendor study with a stated sample size
6	Vendor blog, no sample size
7	Listicle

Source kinds: filed; company-stated; vendor-reported; analyst-derived; measured-by-us.

Case grades:

- **Gold:** holdout, geo-split or switchback (results against a withheld group, region or period).
- **Silver:** before/after with a baseline and an unaffected control.
- **Bronze:** visibility change only.

As graded vs strict rule: strict caps a case at Bronze when any of seven evidence items is missing.

Between 2026-02 and 2026-08 ChatGPT ads went from a US pilot to a stated \$1 billion run rate, while filings record lost clicks

Date	Event	Source kind, tier
2026-02-04	Anthropic: "Claude will remain ad-free"	company-stated, 3
2026-02-09	ChatGPT Ads live, US pilot	company-stated, 3
2026-02-25	NerdWallet cards revenue "decreased 24%", citing "AI overviews and LLMs" (large language models)	filed, 2
2026-08-31	EU designates ChatGPT a VLOSE (Very Large Online Search Engine), 159.1M EU recipients	filed, 2
2026-08-31	ChatGPT Ads "\$1 billion... run rate", "over 40 countries"	company-stated, 3
2026-09-07	Replication: GEO levers move citation on "none" of ten engine families	analyst-derived, 4
2026-09-17	Court exhibit: Microsoft "83-93% drops in click-through rates"	filed, 2

A selection from 54 dated events on file, 26 of them at tier 2. Run rate: revenue annualised from a recent period, not booked revenue.


No sub-market has a measured size: 18 published figures, 15 forecasts, 0 measured sizes

Sub-market	Floor from disclosed figures	Published forecasts
Organic	"\$42.2M–\$48.2M + €2.2M ARR" (annual recurring revenue), 4 vendors, 2025-11 to 2026-05; filed tier 2, company-stated tier 5. A second floor exists (appendix).	2034: \$17.15B–\$32.92B; spread 1.92×; tier 6
Paid	"≥ \$1B run rate", ChatGPT only; company-stated, tier 3, 2026-08-31	2030: ~\$5B–\$100B+; spread ~20×; scopes differ; analyst-derived
Agentic	unknown — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon, 2026-09-23	2030: \$190B–\$5T; spread 26.3×; tier 5–6. Another spread exists (next slide)

Count of 18 figures from the risk register, 2026-09-23 (analyst-derived, tier 6). A later, undated count adds 10 new forecasts. A floor is revenue already disclosed; a forecast is someone's projection. They are never read as each other.

Forecasters disagree by 1.92× on organic, ~20× on paid, and 26.3× to roughly 35× on agentic

Forecast spread, × (bar length proportional; axis starts at 0)



Read: the wider the spread, the less the forecasts agree.

Paid: forecasts differ in scope, so the ~20× mixes definitions.

Agentic: two spreads sit side by side, 26.3× for 2030 only and "roughly 35×" over 2029–30. Neither is dropped.

Tier: organic forecasts tier 6; agentic tier 5–6.

Four engines sell ads inside answers, and none of eight publishes a rate card

Sellers (2026-09-23, tier 3): ChatGPT, Google, Microsoft Copilot, Amazon Rufus (Amazon's shopping assistant). An earlier, undated count reads 5 of 8.

Leaving: Perplexity is "winding down" ads (analyst-derived, tier 5, 2026-02-18).

Price: "0 of 8 publish a rate card" (a published price list; tier 3, 2026-09-22). The only engine price is OpenAI's bid guidance, "\$3–\$5 USD per click" (tier 3, 2026-09-22).

How many ChatGPT advertisers? Sources disagree:

"tens of thousands"
OpenAI, company-stated, tier 3, 2026-08-31. Earlier OpenAI figure: "over 600", 2026-03-26.

820 to 7,378
Advertiser panels, vendor-reported, tier 5, 2026-07.

Four assistants now run checkout, and none of six open checkout protocols states its fee

Live programmes: ChatGPT, Google, Copilot, Perplexity (company-stated, tier 3, 2026-09).

Fees stated:

- Copilot "does not take a commission".
- OpenAI: "Merchants pay a small fee"; rate not disclosed (tier 3, 2025-09-29).


Fee clauses in open checkout protocols (published specifications for assistant-run checkout): 6 of 6 unknown, checked 2026-09-22.

What is not known:

- Checkout volume run by assistants: unknown — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon, 2026-09-23.
- Copilot's fee is carried two ways: 0% in one compilation, unknown in the protocols table (appendix).

Of 34 visibility-tool vendors, none discloses the prompt set behind its score

Vendors disclosing each item, count (2026-09)




Prompt set: the fixed list of questions a vendor runs to score a brand's visibility.

Filed revenue: Semrush, "AI products surpassed \$38 million in ARR" (filed, tier 2, 2025-12). Semrush merged into Adobe 2026-04-28.

Consolidating: Yext bought GoShine 2026-09-01; Sitecore bought Scrunch, announced 2026-06-03.

Spend shows in 8 of 27 buyer segments under both readings; willingness to pay is unknown in all

27 segments by read, count, 2026-09-23 (axis starts at 0)



Segment: sub-market × vertical × buyer size (small business under 100 staff; mid-market 100–999; enterprise 1000 or more).

Reads: spend = money moved; attention = interest only; none = checked and empty.

Strict vs loose: strict leaves a segment unchecked when a channel was blocked; loose counts a blocked-and-logged channel as checked. Earlier readings exist (appendix).

Local multi-location, nine further segments: spend 1, attention 3, none 5 (2026-09-23). Willingness to pay unknown there too.

Most spend signals are job postings, which show headcount, not a category budget

By sub-market (2026-09-23): organic 6 spend segments, paid 1, agentic 1.

By evidence (as read 2026-09-22): six of the seven spend reads rested on job postings. Exactly 1 of 9 small-business segments read spend.

Strongest signal is a ceiling: KKH (a German health insurer) agency framework, "11,000,000.00 EUR", GEO among its themes (filed, tier 2, pulled 2026-09-23). A ceiling for the whole framework, not GEO spend.

What the postings look like:

e.l.f. Beauty: an in-house "AEO ... and GEO ... team"

An agentic-commerce role at "\$110,000.00/yr - \$140,000.00/yr"

Both company-stated, tier 3, 2026-09-22.

No published case uses a holdout, geo-split or switchback, so none reaches Gold

0

cases with a holdout, geo-split or switchback, out of ~980 screened 2026-09-22 and ~1,640 more screened 2026-09-23.

Grade	As graded	Strict rule
Gold	0	0
Silver, 2026-09-22	7	1
Silver added 2026-09-23	4	1


Two Silvers cut the other way: two of the seven Silvers from 2026-09-22, the highest-tier ones, run against the category: NerdWallet's filing (tier 2) and the null replication (tier 4).

Brands that moved report outcomes by observation; at most 6 of 108 used an experiment


Cases by evidence class, count, both readings 2026-09-23 (bars proportional; axis starts at 0)

■ B = Reading B: vendor before/after counted observational A = Reading A: before/after with a control counted experimental

108 cases of brands that moved




52 graded cases



Movers mostly changed content and tooling; 2 of 99 changed paid media

Cases naming each change, count, 2026-09-23 (axis starts at 0)



Base: of the 108 cases of brands that moved, 99 name a change.

Read: movers changed their own content and bought tools; paid media is rare among named changes.

The strongest-tier evidence runs against the category: filings record lost traffic and revenue

Who	What they report	Source kind, tier, date
Microsoft (court exhibit)	"83-93% drops in click-through rates"	filed, 2, 2026-09-17
People Inc (publisher)	"22% decline in Core Sessions", citing Google AI Overviews	filed, 2, 2026-08-03
NerdWallet	Cards revenue "decreased 24%", citing "AI overviews and LLMs"	filed, 2, 2026-02-25
HubSpot (blog)	"lost 5 million visits in a thirty-day period"	company-stated, 3, 2026-09-22
Academic replication	GEO levers move citation on "none" of ten engine families	analyst-derived, 4, 2026-09-07

Read: when the answer is given inside the assistant, fewer users click through. This is risk rank 1 (slide 18).

Brands rarely confirm vendors' success claims, and no paid-ad result is controlled

1 of 168

brands confirm a vendor's claim on their own site (cumulative, 2026-09-23). Earlier count: 0 of 59, 2026-09-22.

0 of 15

paid-ad results in AI answers used a control (2026-09-23).

1.30%

CTR (click-through rate) on 97,000+ impressions, "very few sign-ups" (vendor-reported, tier 5, 2026-08-10).

Read: published success stories are the vendor's account. The brand rarely repeats it, and no published paid-ad result used a control.

The best-evidenced risk is referral collapse: the answer replaces the click

Rank	Risk	Strongest evidence	Tier
1	Referral collapse	"83-93% drops in click-through rates"	2
2	Regulation: disclosure, category exclusions, fair ranking	EU AI Act Art. 50 transparency duties in force 2026-08-02; ChatGPT ad-repository duty 2027-01; UK CMA (Competition and Markets Authority) fair-ranking order 2026-06-17	2
3	Vendor consolidation	Semrush into Adobe; Yext buys GoShine 2026-09-01	2
4	Platform capture: engines sell ads, checkout, measurement themselves	OpenAI Ads Manager, Pixel, Conversions API (application programming interface)	3
7	"Do nothing" suffices	Google: "no additional requirements"; null replication	3; 4

Ranks 1–4 and 7 of eleven risks, ranked by evidence tier, then breadth, not severity. Rank 5, no payback evidence, is slides 13–17. Likelihood, impact, mitigation, owner: no source states one.

The evidence stops at sales lift, prices paid, and what engines actually answer

Not shown by any source:

- Causal sales lift from AI visibility.
- Any sub-market's measured size, price paid, engine rate card, or checkout volume run by assistants.
- What engines answer to buying questions: our own sampling ran one partial day (2026-09-22), then stopped.
- Anything behind blocked or paywalled channels: Indeed, live Reddit, EMARKETER, Crunchbase, Gartner, Forrester (2026-09-23).

Caveats:

- Published cases are winners; full-page re-grading moved 12 cases up, 0 down.
- Figures are unaudited self-reports.
- Load-bearing claims at tier 3 or better: 68.0% (17 of 25, 2026-09-22); 44.4% (12 of 27) and 44.0% (11 of 25), 2026-09-23. Different lists; all stand.
- Absence is only as strong as the channels reached.

Brands are moving and engines are selling, but payback is unproven; four readings are yours to set

#	Reading you choose	What each option gives
1	Silver count: as graded, or strict rule	7 Silver as graded vs 1 strict (2026-09-22); +4 vs +1 (2026-09-23)
2	Demand: strict or loose reading	Spend 8 of 27 either way; none 17 vs 18; not yet checked 1 vs 0 (2026-09-23)
3	Vendor before/after with a control: experimental or observational	Experimental cases among 108 movers: 6 vs 1 (2026-09-23)
4	Which sub-market a hypothetical product would serve	Every table splits on it; spend segments: organic 6, paid 1, agentic 1

What the evidence supports: demand signals and supply are real. **Where it stops:** no controlled proof that it pays. The verdict is yours; no lean is attached.

Appendix A1. Figures carried two ways (1 of 2): both readings stand, neither is dropped

Figure	Reading A, date	Reading B, date
Organic floor	"\$5.3M–\$35.2M annualised", price × count, tier 5, 2026-09	"\$42.2M–\$48.2M + €2.2M ARR", tier 2/5, 2025-11 to 2026-05
Segment tally (spend/attention/none/blank)	7/1/11/8, 2026-09-22; strict 8/1/10/8, 2026-09-23	strict 8/1/17/1; loose 8/1/18/0, 2026-09-23
Engines selling ads	"5 of 8", date not stated	"Four engines", 2026-09-23
ChatGPT advertisers	"over 600", 2026-03-26, tier 5	"tens of thousands", 2026-08-31, tier 3
Silver count	7 as graded, 1 strict, 2026-09-22	+4 as graded, +1 strict, 2026-09-23
Evidence classes, 52 cases	Reading A 10/29/13, 2026-09-23	Reading B 5/34/13, 2026-09-23
Evidence classes, 108 cases	Reading A 6/61/37/4, 2026-09-23	Reading B 1/66/37/4, 2026-09-23

Appendix A1. Figures carried two ways (2 of 2)

Figure	Reading A, date	Reading B, date
Brand-side confirmation	0 of 59, 2026-09-22	1 of 168 cumulative, 2026-09-23
Forecast count	"15 of 18 published figures are forecasts", date not stated	"10 new forecasts added", date not stated
Agentic forecast spread	"roughly 35×", 2029–30	26.3×, 2030 only
Copilot checkout fee	0% in a compilation, date not stated	unknown in the protocols table, date not stated
Share of load-bearing claims at tier 3 or better	"17 of 25 = 68.0%", 2026-09-22	"12 of 27 = 44.4%"; "11 of 25 = 44.0%", 2026-09-23
Scrunch (tool vendor) sale value	"not disclosed", tier 3, 2026-06-03	"\$225 Million", tier 5, date not stated

Reading A counts before/after with a control as experimental; Reading B counts vendor before/after as observational.

Appendix A2. Glossary (1 of 2)

Term	Meaning
AI	Artificial intelligence; here, assistants such as ChatGPT, Gemini, Perplexity, Copilot, Claude
GEO / AEO	Generative / answer engine optimisation: getting a brand named in AI answers, unpaid
SEO	Search engine optimisation, the classical discipline
Copilot, Rufus	Microsoft's assistant; Amazon's shopping assistant
EU, UK, US	European Union, United Kingdom, United States
VLOSE	Very Large Online Search Engine: EU Digital Services Act designation bringing ad-repository duties
EU AI Act Art. 50	Transparency duties for AI interaction and synthetic content
UK CMA	UK Competition and Markets Authority; its 2026-06-17 order requires fair ranking in search AI features
KKH	Kaufmännische Krankenkasse, a German health insurer
ARR; run rate	Annual recurring revenue; revenue annualised from a recent period, not booked
LLM; API	Large language model; application programming interface
Reading A / B	A counts before/after with a control as experimental; B counts vendor before/after as observational

Appendix A2. Glossary (2 of 2)

Term	Meaning
USD, EUR; CTR	US dollar, euro; click-through rate
B2B SaaS	Business-to-business software sold by subscription
High-CPA regulated	Cards, insurance, supplements; CPA is cost per acquisition
Buyer sizes	Small business under 100 staff; mid-market 100–999; enterprise 1000 or more
Spend / attention / none / blank	Segment reads: money moved; interest only; checked and empty; not yet checked
Strict / loose	Strict: a blocked channel leaves a segment unchecked; loose: a blocked-and-logged channel counts as checked
Gold / Silver / Bronze	Case grades: holdout, geo-split or switchback; before/after with baseline and unaffected control; visibility change only
Holdout, geo-split, switchback	Designs that withhold treatment from a comparable group, region, or period
Rate card; prompt set	Published ad price list; fixed list of questions a vendor runs to score visibility
Open checkout protocol	A published specification for assistant-run checkout
EMARKETER	A market-research firm; its reports are paywalled