

Brand visibility inside AI assistants: what the evidence supports

Is there real demand, per buyer segment, for brands to be seen and recommended inside ChatGPT, Gemini, Perplexity, Copilot and Claude? And does it work?

Evidence as of 2026-09

Research only: evidence, not a verdict

Eight parts, in the order a director asks the questions

1. The answer on one page
2. What this market is, and how to read the evidence
3. Why now
4. How big
5. Who is already there: vendors, engines, buyers
6. What the proof shows
7. What could kill it, and where the evidence stops
8. What this leaves to you

No decision is requested. The last slide lists choices only you can make, each with the evidence on either side. Appendix: glossary, figures that carry two readings, anticipated questions.

Brands have started paying to be seen inside AI assistants; the evidence stops before sales

Demand is real, thin and mostly large-company. One engine sells ads at scale. No published case shows sales caused.

Demand	Money spent in 8 of 27 buyer segments (sub-market × industry × company size), on both readings (analyst-derived, 2026-09-23). Price any buyer paid: unknown in all 27.
Supply	ChatGPT Ads: "\$1 billion in annualized revenue run rate" (company-stated, 2026-08-31). Engines: "0 of 8 publish a rate card" (pulled 2026-09-22).
Size	No sub-market has a measured size; every dollar size is a forecast.
Proof	0 causal-design cases in ~980 candidates (2026-09-22) and ~1,640 more (2026-09-23).
Movers	Of 108 cases: content changes 63, new tools 58, paid media 2 (measured-by-us, 2026-09-23).

Three sub-markets, three different buyers, three metrics that are never crossed

Brands want to be named when someone asks an assistant what to buy. They can try three ways.

Organic recommendation

Being cited unpaid in the answer. Also called GEO or AEO (generative / answer engine optimisation).
Buyer: SEO (search engine optimisation) and brand teams.

Paid placement

Ads shown inside or beside the assistant's answer.
Buyer: media buyers.

Agentic commerce

The assistant completes the checkout for the shopper.
Buyer: e-commerce teams.

Visibility: named?

≠

Traffic: clicked?

≠

Sales: bought?

Each is measured separately. Visibility never proves sales.

Every figure carries its source kind; lower tier means stronger evidence

Tier	Source kind	How it is used
1	Measured by us	Highest
2	Filed: annual report, court exhibit, funding filing	High
3	Platform's own pages, docs, pricing	Reliable on existence, biased on framing
4	Panel, clickstream, telemetry	Usable if method published
5	Vendor or agency study with sample, dates, method	Usable, bias flagged
6	Vendor blog, no sample	Marketing, not a source
7	Listicle, roundup, affiliate comparison	Not pulled

Source labels on each figure

measured-by-us · filed · company-stated · vendor-reported · analyst-derived (a forecaster's or survey's figure)

Grades for "it worked" cases

Gold: causal design: holdout, geo-split or switchback test.

Silver: before/after with an unaffected control, not randomised.

Below Silver: no control; cannot show cause.

In 2026 ads went live inside answers while publishers filed traffic losses

The record holds 54 dated events (as of 2026-09-23). Eight carry the story.

Date	Event	Source
2026-02-04	Anthropic: "Claude will remain ad-free"	company-stated
2026-02-09	ChatGPT Ads live, US pilot	company-stated
2026-02-25	NerdWallet credit-card revenue "decreased 24%", citing "AI overviews and LLMs"	filed
2026-05-04	People Inc: "63% decline in Google Search referrals over two years"	filed
2026-06-17	UK competition regulator (CMA): fair-ranking duty covers "search generative AI features"	filed
2026-08-31	EU designates ChatGPT a very large online search engine: 159.1M EU recipients; duties from 2027-01	filed
2026-09-15	Profound, a visibility vendor: "\$180M Series D at \$1.8B valuation"	vendor-reported
2026-09-17	US court exhibit: Microsoft "83-93% drops in click-through rates"	filed

Sources: filed, 2026-02 to 2026-09; company-stated, 2026-02; vendor-reported, 2026-09. Event count: record as of 2026-09-23.

The audiences are already billion-scale; each figure is a different metric, so none is summed

ChatGPT

>1B

weekly active users

2026-08-31

Google AI Overviews

2.5B

"monthly active
users"

2026-05-19

Google AI Mode

1B

monthly active users

2026-05-19

Gemini app

950M

monthly active users

2026-07-22

Attention, not spend

76% of "more than 200" brand and agency ad decision-makers name optimising for AI answers their No. 1 area of increased focus (analyst-derived survey, released 2026-09-10). It measures where attention goes, not dollars moved.


No sub-market has a measured size; every dollar size is a forecast

	Organic recommendation	Paid placement	Agentic commerce
Floor: disclosed revenue	\$42.2M–\$48.2M + €2.2M ARR, 4 vendors, 2025-11 to 2026-05 (filed; company-stated)	"\$1 billion" run rate, one engine (company-stated, 2026-08-31)	unknown — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon, 2026-09-23
Forecast spread	2034: \$17.15B–\$32.92B, 1.92×, 4 forecasts	2030: "just over \$5 billion" to "over \$100 billion", ~20×; scopes differ	2030: \$190B–\$5T, 26.3×, 7 forecasts
Published price	14 of 34 vendors publish one: "\$20/mo" to "\$999/month" (vendor-reported, 2026-09)	OpenAI bid guidance "\$3–\$5 USD per click" (company-stated, 2026-09-22)	Microsoft Copilot fee 0%; OpenAI "a small fee" (2025-09-29), rate undisclosed (company-stated)

ARR: annual recurring revenue. Run rate: current revenue pace annualised. Forecasts are analyst-derived, pulled 2026-09-22. Other bases for the floor and spreads exist; see appendix.

Published forecasts disagree by 26.3× on agentic, ~20× on paid, 1.92× on organic

Highest forecast divided by lowest, as each source states it. Linear scale from zero.




How to read it

A long bar means the published forecasts for that sub-market do not agree. Horizons differ (2034 against 2030), so bars compare disagreement, not size.

The two agentic bars come from two different compilations of forecasts; both stand, neither averaged.

Of 34 organic-visibility vendors, 14 publish a price; none discloses the prompts behind its score

Vendors disclosing each item, out of 34. Bars start at zero.



Consolidation has started

43 companies profiled (2026-09-23).

Largest disclosed line: Semrush "AI products surpassed \$38 million in ARR" (filed, 2026-03-02).

Adobe completed its Semrush purchase 2026-04-28; Sitecore announced Scrunch 2026-06-03; Yext filed on GoShine 2026-09-01.

One engine sells ads at scale; no engine publishes a rate card or its checkout fee terms

Ads inside answers


- Four engines sell ads inside answers: ChatGPT, Google, Microsoft Copilot, Amazon Rufus (company-stated, 2026-09-23). Another count: "5 of 8 engines carry a live product".
- ChatGPT advertisers: "tens of thousands" (company-stated, 2026-08-31) against 820 to 7,378 on outside panels (vendor-reported, 2026-07).
- Price: OpenAI bid guidance "\$3–\$5 USD per click"; "0 of 8" engines publish a rate card.
- Share of answers carrying a ChatGPT ad: 0.8%–26%, panels unreconciled.

Checkout inside the assistant

- Four engines run checkout (company-stated, 2026-09).
- Fees: Copilot 0%; OpenAI "a small fee" (2025-09-29), rate undisclosed.
- Fee clauses in the checkout protocols: "6 of 6" unknown — checked 2026-09-22.

Money is moving in 8 of 27 buyer segments, mostly organic and mostly large companies

27 segments by strongest demand signal, two readings.



Segment: sub-market × industry (skincare, B2B software, high-cost-per-acquisition regulated) × company size. Strict: a blocked channel does not count as checked. Loose: it does.

Who pays

- Spend by sub-market: organic 6, paid 1, agentic 1.
- On the 2026-09-22 read: six of seven spend segments rested on job postings; 5 of the 7 were enterprise; 1 of 9 small-business segments read spend.
- Price any buyer paid: unknown in all 27.
- One segment reads spend to one compiler, attention to another; both stand.

No published case shows sales caused; the strongest evidence is a handful of controlled cases


Rung	What the record holds	Source, date
Gold: causal design	0 in ~980 candidates, plus 0 in ~1,640 more	measured-by-us, 2026-09-22 and 2026-09-23
Silver: controlled, not randomised	Strict grading rule: 1 (2026-09-22), plus "one" (2026-09-23). As first graded: 7, plus "Four".	measured-by-us over vendor-reported cases
Silver that ran negative or null	2 of the 7 as first graded: NerdWallet's filed revenue line and a published replication	filed; 2026-09-22
Brand confirms the vendor's case	1 of 168 (earlier count: 0 of 59, 2026-09-22)	measured-by-us, 2026-09-23

Strict grading rule: any missing evidence item caps a case below Silver. Both readings are kept; neither is averaged.

Sources: measured-by-us screening of vendor-reported and filed cases, 2026-09-22 and 2026-09-23.

Of 108 cases where a company changed something, 6 or 1 count as experiments, depending on the reading


108 cases by evidence class. One bar = 108 cases; linear scale.



Reading 1 counts a vendor's before/after with a control as experimental; Reading 2 counts it as observational. Which governs is an open method choice; both are shown.

Companies that moved changed content and bought tools; paid media barely appears

Named changes by type, in 108 cases. Bars start at zero.



How to read it

Of 108 cases, 99 name a change. Bars count named changes per type as the record states them; they are not shares of 108.

The count is of actions named, not of visibility, traffic or sales moved.

The filed results run negative; the positive results are vendor-measured and never per engine

Negative

People Inc: "22% decline in Core Sessions" (filed, 2026-08-03).

HubSpot blog: "5 million" fewer visits in thirty days (company-stated, 2026-09-22).

Null

Optimisation levers move citation on "none" of ten engine families (analyst-derived, 2026-09-07).

Positive, vendor-measured

Visits: "+300% more site visits from AI Search", "February 23 to June 29, 2026".

Conversion: Adobe "AI Conversion Now 42% Higher" than non-AI visits (2026-03); Shopify "nearly 50% higher rates than organic search" (Q1 2026).

Q1: first calendar quarter. Neither conversion figure is broken out per engine; per-engine conversion is unknown — checked Adobe, Salesforce, Shopify, 2026-09-22.


Eleven ranked threats; no source states a likelihood or impact for any of them

Rank	Threat to "marketing inside assistants pays"	Evidence on record
1	Referral collapse	Court exhibit: "83-93% drops in click-through rates"
2	Regulation	Disclosure duties; Google AI Overviews ads exclude finance, healthcare
3	Vendor consolidation	Adobe–Semrush, Sitecore–Scrunch, Yext–GoShine
4	Platform capture	OpenAI: "more than 50 technology and measurement partners"; Google: "you can't opt out"
5	No payback evidence	"0 Gold; 0 of 15 ad results controlled"
6	Measurement opacity	ChatGPT ad presence 0.8%–26% across panels
7	Substitutes	Google: no "special optimizations" are "necessary"
8	Engine policy reversal	Claude ad-free; Perplexity winding down its ads
9	Agentic rails	Engine-owned checkout, fees undisclosed
10	Manipulation	Google spam policy names "manipulate generative AI responses" (2026-08-28)
11	Unsizeable market	"18 figures, 15 forecasts, 0 measured sizes"

Rank order from the compiled risk register, 2026-09-23. Sources: filed; company-stated, 2026-08 to 2026-09; analyst-derived. Likelihood, impact, mitigation, owner: no source states one.

No count reaches the 80% bar: 44.0% to 63.2% of load-bearing claims rest on tier 3 or better

Share of claims at tier 1–3. Bars start at zero; linear scale.



How to read it

Load-bearing claim: a claim a finding rests on. Each list is counted three ways: all claims; without the vendor case corpus; without claims that have no source tier.

Counts as of 2026-09-23. An earlier count, 2026-09-22, read 68.0%.

What the record cannot yet say, and why

Unknown	Status
Price any buyer paid	unknown in all 27 segments
Agentic spend floor	unknown — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon, 2026-09-23
Conversion per engine	unknown — checked Adobe, Salesforce, Shopify, 2026-09-22
Checkout protocol fees	"6 of 6" clauses unknown, 2026-09-22
Hypotheses scored	26 of 32; six need our own prompt runs against the engines, skipped 2026-09-23

Channels the research could not reach

Still blocked, 2026-09-23: reddit.com, indeed.com, emarketer.com, ted.europa.eu, perplexity.ai/hub.

Behind paid credentials: Crunchbase, PitchBook, Gartner, Forrester, EMARKETER.

No decision is requested; five choices are yours, each with evidence on both sides

Choice you own	Evidence that bears on it
Which sub-market any reading rests on	Organic: 6 spend segments, no measured size. Paid: 1 spend segment, one engine at "\$1 billion" run rate. Agentic: 1 spend segment, floor unknown.
Which reading governs where two stand	Silver cases: 7 as first graded, 1 under the strict rule. Experimental cases: 6 or 1. Segments checked with no demand: 17 strict, 18 loose.
Whether our own engine sampling reopens	Six of 32 hypotheses need it; ad-presence panels disagree, 0.8%–26%.
Whether paid-credential sources enter scope	Crunchbase, PitchBook, Gartner, Forrester, EMARKETER hold funding and dollar-proxy data the record lacks.
Whether one question returns to scope	"which of the gaps above, if any, is reachable by a team whose cost advantage is engineering rather than distribution?"

Glossary

Term	Meaning
AI; assistant; engine	Artificial intelligence; one assistant product (ChatGPT, Gemini, Perplexity, Copilot, Claude, Amazon Rufus)
AI Overviews; AI Mode	Google's AI answer block on a results page; Google's conversational search tab
GEO; AEO	Generative engine optimisation; answer engine optimisation: names for the organic sub-market
SEO	Search engine optimisation
B2B software	Business-to-business software sold as a service
High-cost-per-acquisition regulated	Industries where each customer is expensive to win: credit cards, insurance, supplements
Segment (cell)	One sub-market × industry × company size; read as spend (money moved), attention (interest only), none — checked, or blank (not yet checked)
Strict; loose	Whether a blocked-and-logged channel counts as checked for a "none" read: no; yes
Small business; enterprise	Buyer sizes; mid-market sits between

Glossary, continued

Term	Meaning
Tier	Source provenance score; lower is stronger (scale on the "How to read the evidence" slide)
Load-bearing claim	A compiled claim a finding rests on
Gold; Silver	Case grades: causal design (holdout, geo-split, switchback); before/after with an unaffected control
As first graded; strict rule	Grade as the screener assigned it; grade when any missing evidence item caps a case below Silver
Reading 1; Reading 2	Vendor before/after with control counted experimental; counted observational
ARR; run rate	Annual recurring revenue; current revenue pace annualised
MAU; LLM	Monthly active users; large language model
CMA	UK Competition and Markets Authority
Very large online search engine	EU Digital Services Act designation carrying extra duties
US; UK; EU; USD	United States; United Kingdom; European Union; US dollars
Own engine sampling	The research's own prompt runs against the engines

Where a figure has two readings or two dates, both stand

Figure in the deck	Other figure	Why two
Segments strict 8/1/17/1, loose 8/1/18/0 (2026-09-23)	7/1/11/8, first read (2026-09-22); strict 8/1/10/8, loose 8/1/18/0, intermediate read	dates, readings
Silver, strict rule: 1, plus "one"	As first graded: 7 (2026-09-22), plus "Four" (2026-09-23)	two readings
Classes, Reading 1: 6/61/37 (2026-09-23)	Reading 2: 1/66/37, same date	two readings
Brand confirms case: 1 of 168 (2026-09-23)	0 of 59 (2026-09-22)	two dates
Four engines sell ads (2026-09-23)	"5 of 8 engines carry a live product"	two bases
ChatGPT Ads "\$1 billion" run rate (2026-08-31)	"\$100 million annualized" (2026-03-26)	two dates
Advertisers "tens of thousands" (2026-08-31)	"over 600" (2026-03-26)	two dates
Organic floor \$42.2M–\$48.2M + €2.2M ARR	\$5.3M–\$35.2M annualised, price × vendor count (2026-09)	two bases

Figures carried twice, continued

Figure in the deck	Other figure	Why two
Agentic spread 26.3×, 2030	"roughly 35×", 2029–2030	two bases
Copilot checkout fee 0%	"1 of 4" checkout programmes disclose a fee (2026-09)	two reads
ChatGPT ad presence 0.8%–26%	"26%" one panel; "4.47%" US, 2026-03 to 2026-05; "0.00 percent" of 169,560 UK scrapes, 2026-06	panels, unreconciled
Tier-3-or-better share, six counts (2026-09-23)	17 of 25 = 68.0% (2026-09-22)	two dates
Skincare agentic enterprise segment: spend	attention, per a second compiler	two compilers
Hypotheses 26 of 32 scored (2026-09-23)	15 of 23 (2026-09-22)	two dates
Peec AI ARR \$4M (2025-11)	\$10M (2026-05)	two dates
Profound round: vendor-reported	company-stated in an earlier brief of 2026-09-23	two labels

Questions a director is likely to ask, answered from the record

Is there a market today? Money moves in 8 of 27 segments. No sub-market has a measured size.

How large could it get? Only forecasts exist; for agentic, 2030 forecasts run \$190B–\$5T, a 26.3× spread.

Does it work? 0 causal-design cases; brands confirm a vendor's case in 1 of 168.

Who pays? Mostly enterprise: 5 of the 7 spend segments on the 2026-09-22 read; 1 of 9 small-business segments.

What do buyers pay? Unknown in all 27 segments. Vendor list prices run "\$20/mo" to "\$999/month".

Why does traffic matter here? A court exhibit shows "83-93% drops in click-through rates"; People Inc filed a "63% decline in Google Search referrals over two years".

Where is an entrant most exposed? Platform capture: engines sell ads, checkout and measurement themselves, with "more than 50 technology and measurement partners" at OpenAI.

What would change the reading? One controlled result on sales; the six unscored hypotheses; the paid-credential sources.

How the evidence behind this deck was built

- Internet sources only: company filings, court exhibits, regulator designations, platform help and pricing pages, panel vendors, vendor case studies, analyst forecasts, job postings.
- No interviews, no outreach. Our own prompt runs against the engines were skipped (2026-09-23).
- Evidence frozen at one snapshot; nothing later is cited. Every figure traces to a verbatim source pull held in the research record, with line references kept outside this deck.
- Conflicting figures sit side by side; none is averaged. Where the record has no figure, the deck says "unknown" and names the channels checked.
- The research produces evidence. The verdict belongs to the reader.