

BRIEFING FOR THE DIRECTOR

Brand visibility and product recommendation inside AI assistants: what the evidence shows

Is there real demand, per buyer segment, to be named or recommended when people ask ChatGPT, Gemini, Perplexity, Copilot or Claude what to buy, and does anything show it works?

Research only: evidence, not a verdict

Evidence as of 2026-09

AI: artificial intelligence. Every acronym is defined on first use or in the glossary, last appendix slide.

The deck walks a newcomer through seven steps, in order

1. **What this is:** three markets, and how to read the evidence
2. **Why now:** ads, regulation and referral loss landed in 2026
3. **How big:** what is disclosed, what is forecast, what is unknown
4. **Who is already there:** engines, vendors, buyers
5. **Does it work:** how strong the published proof is
6. **What could kill it:** the top five risks and the evidence against each
7. **What is asked of you:** no decision; the reading choices that move the figures

Appendix: figures with two readings, risks ranked 6 to 11, blocked channels, sources and caveats, glossary.

Sellers are ahead of buyers, proof trails both, and no sub-market has a measured size

- **Buyers move, thinly:** 8 of 27 buyer segments show spend (strict reading, 2026-09-23), mostly from job postings; willingness to pay is unknown in all 27.
- **Sellers are ahead:** four engines sell ads inside AI answers (2026-09-23); ChatGPT ads reached a "\$1 billion in annualized revenue run rate" (2026-08-31).
- **Proof trails:** no published case uses a causal design; every positive case with a control is a vendor or agency study; the filed cases run negative.
- **Size is unmeasured:** of 18 published size figures, 15 are forecasts and 3 measured, none a sub-market size.

What the evidence supports: activity on both sides is documented. **Where it stops:** no buyer price paid, no rate card, no sub-market size, no measured sales lift. The verdict is yours.

The market splits three ways, and sales are never inferred from visibility

Brands want to be named when someone asks an AI assistant what to buy. Three ways to get there:

Organic recommendation

Unpaid appearance in the answer.
Served by software tools and agencies.
Also called GEO (generative engine optimisation) or AEO (answer engine optimisation).

Paid placement

Ads shown inside the AI answer, sold by the engine that runs the assistant.

Agentic commerce

The assistant itself runs the checkout; the merchant is paid through it.

Three results are measured apart and never crossed: **visibility** (named in the answer), **traffic** (clicks to the brand) and **sales**. A change in visibility alone is never read as proof of sales.

Every figure carries a tier from 1, strongest, to 7, weakest; published cases carry a grade

Tier	What sits there
1	"Highest. Cite freely"
2	Filed: regulatory, court or company filings
3	Platform primary page: the company's own statement
4	Panel, clickstream or telemetry data
5	Vendor or agency study
6	"Marketing. Not a source"
7	"Do not pull"


Case grade	Design
Gold	Causal: holdout, geo-split or switchback
Silver	Before/after with an unaffected control
Bronze	Visibility change only

Source kinds on every footer: filed, company-stated, vendor-reported, analyst-derived, measured-by-us (counted by this research).

Ads, regulation and referral loss all landed between 2026-02 and 2026-09

2026-02-09 ChatGPT ads go live, US pilot. <i>Company-stated, tier 3</i>	2026-02-25 NerdWallet filing: cards revenue "decreased 24%", citing "AI overviews and LLMs" (large language models). <i>Filed, tier 2</i>	2026-06-17 UK CMA (Competition and Markets Authority) fair-ranking requirement covers "search generative AI features". <i>Filed, tier 2</i>
2026-08-02 EU AI Act Article 50 in force. <i>Filed, tier 2</i>	2026-08-31 ChatGPT ads at "\$1 billion" run rate, "over 40 countries" (company-stated, 3). EU designates ChatGPT a very large search engine: 159.1M EU recipients, ad-repository duty from 2027-01 (filed, 2).	2026-09-17 Court exhibit: Microsoft "83-93% drops in click-through rates". <i>Filed, tier 2</i>

Full timeline: 54 events, 2023-04-25 to 2027-01, by evidence tier (events, linear)



Of 18 published size figures, 15 are forecasts and 3 measured; none sizes a sub-market


Sub-market	Floor: disclosed revenue only	Forecast spread	Prices published
Organic	\$42.2M–\$48.2M + €2.2M ARR (annual recurring revenue), 4 vendors, 2025-11 to 2026-05 (filed 2; company-stated 5)	2034: \$17.15B–\$32.92B, 4 forecasts, tier 6	14 of 34 vendors publish a price, "\$20/mo" to "\$999/month" (vendor-reported, 2026-09)
Paid	≥ \$1B run rate, one engine, 2026-08-31 (company-stated, 3)	2030: "just over \$5 billion" to "over \$100 billion", ~20×, tier 5	0 of 8 engines publish a rate card; OpenAI "\$3–\$5 USD per click" (company-stated, 3)
Agentic	unknown — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon 2026-09-23	2030: \$190B–\$5T, 26.3×, 7 forecasts, tier 5–6	Copilot "does not take a commission"; ChatGPT "a small fee"; 6 of 6 protocol fee clauses unknown — checked 2026-09-22

Floor and spread each have a second reading; both stand side by side in the appendix. Run rate: current revenue annualised.

Forecasts for the same year disagree by ~20× for paid and 26.3× for agentic

Organic, 2034

4 forecasts · tier 6




Paid, 2030

spread ~20× · tier 5



Agentic, 2030

7 forecasts · 26.3× · tier 5–6



Each row has its own linear axis in US\$ billions, starting at zero. Rows are not comparable to each other.
Second reading of the agentic spread: "roughly 35×" over 2029–2030 (2026-09-22). These are forecasts, not sizes.

Assistant use is large, but AI's share of search is small and AI referral traffic is still growing

">1 billion"

ChatGPT weekly active users

Company-stated, tier 3, 2026-08-31

1.65%

AI tools' share of US desktop search events, Jan–Mar 2026

Vendor-reported, tier 5

770.7M

AI referral visits per month on average, "+117.4%" year on year, Jun 2025–May 2026

Vendor-reported, tier 4

Three different metrics from three sources; they are shown apart and not combined. Together they bound the audience, not the market.

Engines own the inventory: four sell ads in answers, none publishes a rate card

- **Selling ads, as of 2026-09-23:** ChatGPT, Google, Microsoft Copilot, Amazon Rufus (company-stated, 3). An earlier count reads 5 of 8.
- **Holding out:** "Claude will remain ad-free" (Anthropic, 2026-02-04, tier 3).
- **Price:** 0 of 8 engines publish a rate card; OpenAI states "\$3–\$5 USD per click" (tier 3).


How many advertisers buy ChatGPT ads? Two answers, not averaged

Figure	Who says
"tens of thousands of advertisers"	OpenAI, 2026-08-31, company-stated, tier 3
820 to 7,378	Third-party panels, vendor-reported, tier 5

1 of 8 engines discloses ad revenue: OpenAI.

How often ChatGPT shows an ad depends on who measured and where: 26% to 0.00%

Share of sampled ChatGPT answers or queries carrying an ad (% , linear, 0–30)




Same metric, different samples and dates, all tier 5. Shown side by side, not averaged.

Checkout inside the assistant is live at 4 assistants; what it costs the merchant is mostly undisclosed

- **Live:** 4 assistants run checkout (company-stated, tier 3, 2026-09).
- **Fee statements:** 1 of 4 programmes states its fee position: Copilot "does not take a commission". ChatGPT charges "a small fee on completed purchases", rate undisclosed.
- **Protocols:** 6 of 6 protocol fee clauses unknown — checked 2026-09-22.
- **Merchant control:** Shopify merchants were auto-enrolled into Copilot Checkout after an opt-out window (tier 3).
- **Size:** no disclosed revenue floor — checked OpenAI, Shopify, PayPal, Stripe, Google, Amazon 2026-09-23.

Organic-side vendors disclose little, and incumbents are already buying them

What 34 rostered vendors disclose (vendors, 0–34)




n: sample size. Prompt-set count tier 3.

- **Semrush** filed "AI products surpassed \$38 million in ARR" (2025-12, filed 2); Adobe bought it, 2026-04-28 (filed 2).
- **Scrunch to Sitecore**, 2026-06-03: value "not disclosed" (company-stated, 3) beside "\$225 Million" in the press (analyst-derived, 5).
- **Profound**: "\$1.8B valuation" (tier 3).
- **Cases**: of brands named in vendor case studies, 1 of 168 corroborates on its own site (measured-by-us, 2026-09-23).

8 of 27 buyer segments show spend; the rest depends on how a blocked channel counts

Segment = sub-market × vertical (skincare and beauty; business software; high-CPA regulated, where winning a customer costs much) × company size. CPA: cost per acquisition.

Buyer segments by strongest demand signal, read 2026-09-23 (segments, linear, 0–27)



Strict: every signal channel must be run, including a still-blocked one. Loose: a blocked-and-logged channel counts as checked.

Spend by sub-market: organic 6, paid 1, agentic 1. **Basis:** on 2026-09-22, six of the seven spend reads rested on job postings (tier 3). **Willingness to pay:** unknown in all 27.

Named buyers are moving through hires, filings and tenders, not disclosed budgets

Buyer	What the record shows	Kind, tier
e.l.f. Beauty	AEO/GEO team; Google "Sponsored deal" pilot; agentic role at "\$110,000.00/yr - \$140,000.00/yr"	Job posting, tier 3
Coty	Filing names "generative engine optimization", 2026-08-20	Filed, tier 2
Cigna, Pennylane	Named movers in the buyer read	Analyst-derived
KKH, German statutory health insurer	Agency framework naming GEO, ceiling "11,000,000.00 EUR"	Filed, tier 2

Local chains (multi-location businesses), read separately 2026-09-23: spend 1, attention 3, none 5.

No screened case reaches a causal design; every positive controlled case is a vendor study


Rung	Screen of 2026-09-22	Added 2026-09-23
Published cases screened	~980	~1,640 more
Gold (causal)	0	0
Silver (control), tier 5	7 raw, 1 rule-1	+4 raw, +1 rule-1
Silver, negative or null	2 of 7	+4

Raw: grade as first given. **Rule-1:** literal rule; missing any of seven items (brand, engine, date window, baseline, intervention, sample size, who measured) caps a case at Bronze. Counts measured-by-us.

- **Only one Silver reaches sales:** NerdWallet's filing, a decline attributed by management, not measured lift.
- **Negative replication:** levers moved citation on "none" of ten engine families (tier 4).
- **Paid ads:** 15 results screened, 0 with a control (tier 5).
- **Re-grading** on full pages moved 12 cases up and 0 down.

Most documented cases are observational; experimental cases fall from 10 to 5 under the stricter reading

52 published cases by design, 2026-09-23 (cases, linear, 0–52)




R1 counts a before/after with a control as experimental. R2 counts vendor before/after as observational.

What brands changed: of 108 documented brand moves, 99 name a change; paid media appears in 2, both Google's Direct Offers pilot (a "Sponsored deal" ad unit in Google's AI Mode). Measured-by-us.

No count of the claims this research rests on meets its own bar of 80% at tier 3 or better

Load-bearing claims at tier 3 or better, recount 2026-09-23 (% of claims, linear, 0–100)



Earlier count on the 25-claim list: 68.0% (17 of 25, tiers of 2026-09-22). Case corpus: the vendor case studies. Meta-claims: claims about the research itself. Load-bearing: claims the conclusions rest on.

The top five risks rest on filed or platform evidence; each has evidence against it

Rank	Risk	Strongest evidence (tier)	Against it (tier)
1	The answer replaces the click	"83-93%" CTR (click-through rate) drops; NerdWallet; "22% decline in Core Sessions" (2)	Referral visits "+117.4%" (4)
2	Regulation	AI Act Art. 50; ad-repository duty; CMA order (2)	"No enforcement action found against any engine" (2)
3	Incumbents absorb the tools	Semrush to Adobe; Yext buys GoShine (2)	Profound "\$1.8B valuation" (3)
4	Engines self-serve ads, checkout, measurement	OpenAI Ads Manager, Pixel, Conversions API; Google "you can't opt out" (3)	0 native visibility tools at OpenAI, Anthropic (3)
5	No payback evidence	HubSpot blog "5 million" fewer visits (3); ad test CTR 1.30%, "very few sign-ups" (5)	Citations 48 vs 426, one experiment (5)

Register of 2026-09-23, ranked by evidence tier. Likelihood, impact, mitigation: no source states one. Ranks 6 to 11 in the appendix.

No decision is requested; three open reading choices move the headline figures

Reading choice (yours to make)	Option A	Option B
Silver cases: as graded, or under the literal rule?	As graded: 7	Literal rule: 1
Does a blocked-and-logged channel count as checked?	Strict: 26 of 27 segments closed	Loose: 27 of 27
Is vendor before/after with a control experimental?	R1: 10 of 52	R2: 5 of 52

Not in the record: any price a buyer paid; any engine rate card; any sub-market size; what assistants answer under this research's own measurement (sampling held 2026-09-22, skipped 2026-09-23). Blocked and paywalled channels: appendix.

Appendix

- A · Figures carrying two readings or two dates
- B · Risk register, ranks 6 to 11
- C · Blocked and paywalled channels
- D · Sources and caveats
- E · Glossary

Size and engine figures: both readings stand, nothing averaged

Figure	Reading A	Reading B
Organic floor	\$42.2M–\$48.2M + €2.2M ARR, disclosed, 4 vendors (2026-09-23)	\$5.3M–\$35.2M annualised, price × customer count, 4 of 34 vendors (2026-09)
Peec AI ARR, inside the floor	">\$4 million", 2025-11	"10m", 2026-05
Size figures	18 figures, 15 forecasts, 3 measured	10 new forecasts, all forecast-labelled (2026-09-23)
Agentic forecast spread	26.3×, 2030 only (2026-09-23)	"roughly 35×", 2029–2030 (2026-09-22)
Engines selling ads	four, as of 2026-09-23	5 of 8 (2026-09-22)
ChatGPT advertisers	"tens of thousands", 2026-08-31	820 to 7,378 on panels; "over 600", 2026-03-26

Ad presence, fees, deal value and corroboration: both readings stand

Figure	Reading A	Reading B
ChatGPT ad presence	26% US desktop, 2026-06; 25.94% of 50,006 US prompts; 24.7% US, 2026-07	4.47% of US queries, 2026-03 to 2026-05; ~0.8% of 500+ prompts, launch weeks; 0.00% of 169,560 UK scrapes, 2026-06
Checkout fee statements	1 of 4 programmes states a fee position, Copilot	ChatGPT "a small fee on completed purchases", rate undisclosed
Scrunch sale value	"not disclosed", tier 3	"\$225 Million", tier 5
Brand-side corroboration	0 of 59 (2026-09-22)	1 of 109; cumulative 1 of 168 (2026-09-23)
Skincare × agentic × enterprise	spend (buyer file)	attention (market file)

Buyer segment tallies and case counts moved within two days

Figure	Reading A	Reading B
Segment tally, 27 cells (spend / attention / none / unchecked)	7 / 1 / 11 / 8 (2026-09-22); strict 8 / 1 / 10 / 8, loose 8 / 1 / 18 / 0 (2026-09-23)	strict 8 / 1 / 17 / 1, loose 8 / 1 / 18 / 0 (2026-09-23, later re-probe)
Segments meeting the bar	19 of 27 (2026-09-22); strict 19 of 27, loose 27 of 27 (2026-09-23)	strict 26 of 27, loose 27 of 27 (2026-09-23, later)
Cases screened	~980 (2026-09-22)	~1,640 more (2026-09-23)
Silver count	7 raw, 1 rule-1 (2026-09-22)	+4 raw, +1 rule-1 (2026-09-23)
Negative or null	2 of 7 Silvers (2026-09-22)	+4 graded, +1 ungraded replication (2026-09-23)
Sitefire / Jerry case grade	Bronze, one census	Silver, a later census

Case design, brand moves, claim tiers and own sampling: both readings stand

Figure	Reading A	Reading B
Design, 52 cases	R1 10 / 29 / 13	R2 5 / 34 / 13
Design, 108 cases	R1 6 / 61 / 37 / 4	R2 1 / 66 / 37 / 4
108 cases, what is counted	99 name a change	67 metric-moved, 37 action-only, 4 neither
Tier-3 share, 25-claim list	17 of 25 = 68.0%, 2026-09-22 tiers	11 of 25 = 44.0%; 11 of 18 = 61.1%; 11 of 22 = 50.0% (2026-09-23)
Tier-3 share, 27-claim list	12 of 27 = 44.4%	12 of 19 = 63.2%; 12 of 24 = 50.0%
Own sampling of assistants	held 2026-09-22 22:40	skipped 2026-09-23

Ranks 6 to 8: measurement opacity, "do nothing", engine reversal

Rank	Risk	Strongest evidence (tier)	Against it (tier)
6	Measurement opacity	"0 of 34 rostered vendors disclose a fixed published prompt set with n" (3); Google Search & Other "over \$63 billion", no AI line (3); ad presence 0.8%–26% (5)	OpenAI Pixel, Conversions API (3); Ahrefs "454M+ prompts" partial disclosure (3)
7	"Do nothing" suffices	Google: "no additional requirements... nor other special optimizations necessary" (3); levers move citation on "none" of ten (4)	IAB (Interactive Advertising Bureau) "No. 1 area of increased focus... 76%" (4; attention, not spend)
8	Engine policy reversal	"Claude will remain ad-free", 2026-02-04 (3); Gemini app "not rushing anything here" (3); Perplexity "winding down... by the end of 2026" (5)	"over 40 countries", "\$1 billion" run rate (3); Amazon prompts GA (general availability), Copilot ads live (3)

Register of 2026-09-23. Likelihood, impact, mitigation: no source states one.

Ranks 9 to 11: checkout fees, manipulation, an unsizeable market

Rank	Risk	Strongest evidence (tier)	Against it (tier)
9	Agentic rails: fees undisclosed	Shopify merchants auto-enrolled into Copilot Checkout after an opt-out window (3); 6 of 6 protocol fee clauses unknown (3)	Copilot "does not take a commission or affiliate fee" (3); merchant stays "merchant of record" (3)
10	Manipulation and countermeasures	Social-proof wording δ Rate "+334%" on Claude 3.5 Sonnet (3); Google spam policy names "manipulate generative AI responses", 2026-08-28 (3)	No top-priority engine names corpus seeding or comparison-page farming (3)
11	Unsizeable market	18 figures, 15 forecasts, 0 measured sizes (6); agentic 2030 spread 26.3 $\times$ (5–6)	"AI products surpassed \$38 million in ARR" (2); "\$1 billion" run rate (3)

δ Rate: a paper's measure of change in how often a product is recommended.

Absence of evidence is only as strong as the channels that could be reached

- **Still blocked, re-probed 2026-09-23:** reddit.com, indeed.com, emarketer.com, ted.europa.eu (later reached), perplexity.ai/hub.
- **Needs payment:** crunchbase.com, gartner.com, forrester.com ("individual purchase (\$1495)").
- **What they would unlock:** PitchBook's free view already shows some rounds (AthenaHQ \$2.1M; Peec seed \$8.08M).
- Gartner holds report bodies behind 8 AEO/GEO titles, 2025-12 to 2026-08.
- Forrester sells "The AEO Technologies Landscape, Q3 2026".
- The Scrunch "\$225M" figure is already visible in a Bloomberg URL.

Caveats: internet-only, unaudited, and pulled over two days in a fast-moving category

- **Input:** a frozen research snapshot, read through a machine-checked fact pack and an earlier director brief of 2026-09-23. Every number traces to a logged repository line.
- **No interviews:** no buyer, vendor or engine was interviewed.
- **Unaudited:** prices, customer counts, funding and cases are vendor or company statements.
- **Forecasts** are shown as forecasts, never as sizes.
- **Job postings** show intent, not category spend.
- **Own measurement** of assistant answers is one partial day.
- **Dates:** pulls 2026-09-22 to 2026-09-23; the category turns over fast.

Acronyms and terms, A to G

AEO	Answer engine optimisation, an alias for the organic sub-market
AI	Artificial intelligence
AI Overviews, AI Mode	Google's AI answer block on a results page; Google's conversational search tab
API	Application programming interface
ARR	Annual recurring revenue
Art.	Article of a regulation
CMA	Competition and Markets Authority, UK
CPA	Cost per acquisition; high-CPA verticals pay much to win a customer
CTR	Click-through rate: clicks divided by impressions
Direct Offers	Google's pilot ad unit in AI Mode, labelled "Sponsored deal"
EU, UK, US; EUR, USD	European Union, United Kingdom, United States; euro, US dollar
GA; GEO	General availability; generative engine optimisation (never geographic here)

Acronyms and terms, G to V

Gold, Silver, Bronze	Case grades: causal (holdout, geo-split, switchback); before/after with an unaffected control; visibility change only
IAB	Interactive Advertising Bureau, an ad-industry body
KKH	Kaufmännische Krankenkasse, a German statutory health insurer
LLM; n; Q3	Large language model; sample size; third calendar quarter
Raw, rule-1	Case grade as first given; literal rule, where a missing bar item caps a case at Bronze
R1, R2	R1 counts before/after with a control as experimental; R2 counts vendor before/after as observational
Run rate	Current revenue annualised
Strict, loose	Strict needs every signal channel run; loose counts a blocked-and-logged channel as checked
Tier	Evidence strength, 1 strongest to 7 weakest
VLOSE	Very Large Online Search Engine, an EU Digital Services Act designation carrying ad-repository duties