

EVIDENCE PRESENTATION FOR A DIRECTOR

Brand visibility and product recommendation inside AI assistants: is there real, paying demand, and does it work?

Assistants covered: ChatGPT, Gemini, Perplexity, Copilot, Claude

Research only: evidence, not a verdict

Evidence as of 2026-09

Paid placement inside AI assistants already sells; brand demand is real but thin, enterprise-led and unpriced; no published case shows it raises sales

Supply: engines sell ads

\$1 billion

ChatGPT ads "annualized revenue run rate" (company-stated, 2026-08-31)

Demand: who is paying

8 of 27

buyer segments show spend (measured-by-us, 2026-09-23).
Willingness to pay: unknown in 27 of 27

Proof: does it work

0 causal cases

in ~980 screened (2026-09-22) and ~1,640 more (2026-09-23),
measured-by-us

Size: how big

15 of 18

published size figures are forecasts; no sub-market has a
measured size (analyst-derived, pulled 2026-09-22)

Where the evidence stops: effect on sales, prices buyers paid, and engine rate cards.

WHAT THIS IS

Brands want to be named when a person asks an AI assistant what to buy; the market splits three ways

Organic recommendation

The brand is mentioned in the answer without paying the engine. Brands buy tools and content work to influence it.

Paid placement

Ads shown inside or beside the assistant's answer, sold by the engine.

Agentic commerce

The assistant completes the purchase itself, through a checkout the engine runs.

Three outcomes, always measured separately, never merged: visibility (is the brand named), traffic (does anyone click through), sales (does anyone buy). Evidence for one is not evidence for another.

Every figure carries a source kind; sources are graded by strength, and results by how causal they are

Tier	Source strength (1 strongest, 7 weakest)
1	Measured by us
2	Filed: regulatory filing, court exhibit
3	Platform primary: own docs, pricing page
4	Panel or clickstream data
5	Vendor or agency study with method
6	Vendor marketing: not a source
7	Listicle or roundup: not pulled

Result grades

- **Gold**: holdout, geo-split or switchback test. Causal.
- **Silver**: before/after with baseline and an unaffected control. Correlational.
- **Bronze**: visibility change only.

Source kinds

filed · company-stated · vendor-reported · analyst-derived · measured-by-us · relayed (press report, unconfirmed)

Segment reads

spend (money seen), *attention* (interest only, never counted as spend), *none*, *blank* (not checked).

Across 2026 the engines began selling placement and checkout, and regulators and courts began to act

Date	Event	Source kind
2026-02-04	Anthropic: "Claude will remain ad-free"	company-stated
2026-02-09	ChatGPT ads go live, United States (US) pilot	company-stated
2026-03-25	Amazon sponsored prompts generally available, US	company-stated
2026-05-19	Google AI Mode reaches "1B" monthly active users	company-stated
2026-08-02	European Union (EU) AI Act Article 50 transparency duties in force	filed
2026-08-31	ChatGPT designated a very large online search engine; 159.1M EU recipients	filed
2026-08-31	ChatGPT ads "\$1 billion" run rate; >1B weekly active users	company-stated
2026-09-15	Profound (visibility-tool vendor): "\$180M Series D at \$1.8B valuation"	vendor-reported
2026-09-17	Court record, Microsoft: "83-93% drops in click-through rates"	filed

Sources: company-stated, 2026-02 to 2026-08; filed, 2026-08 and 2026-09; vendor-reported, 2026-09


No sub-market has a measured size: disclosed revenue floors are small, and forecasts disagree widely

Sub-market	Disclosed floor (revenue actually reported)	Forecast (not measured)
Organic	\$5.3M–\$35.2M annualised, from 4 of 34 vendors (vendor-reported, 2026-09) Semrush AI products "\$38 million in ARR" (filed, 2026-03-02)	2034: \$17.15B–\$32.92B across 4 forecasts, all tier 6 (analyst-derived, pulled 2026-09-22)
Paid	ChatGPT "\$1 billion" run rate only (company-stated, 2026-08-31)	2030: EMARKETER, US only, "\$68.25 billion" 2030: WPP Media, global, "over \$100 billion" (analyst-derived, pulled 2026-09-22)
Agentic	None computable: 1 of 4 checkout programmes states a fee (company-stated, 2026-09)	2030: \$190B–\$5T across 7 forecasts, a 26.3× spread; another compile states "roughly 35×" (analyst-derived, pulled 2026-09-22)

Forecast horizons, scopes and definitions differ, so the figures are not placed on one chart. ARR = annual recurring revenue.

Thirty-four tool vendors sell visibility tracking, but few disclose price, funding or revenue, and none a fixed test set

Of 34 rostered vendors, how many disclose... (count of vendors)



Dashed outline = all 34 vendors

Published prices run from "\$20/mo" to "\$999/month".

Consolidation has started

- Adobe completed its Semrush purchase 2026-04-28 at "\$12.00/share cash" (filed).
- Sitecore announced buying Scrunch 2026-06-03. Value undisclosed by Sitecore (company-stated); \$225M per Bloomberg, unconfirmed (relayed). Both stand.
- Profound: "\$180M Series D at \$1.8B valuation" (vendor-reported, 2026-09-15).


The engines themselves sell placement, but none publishes a rate card, and measured ad frequency does not reconcile

Of 8 engines tracked (count of engines)



- Only stated price: OpenAI bid guidance, "\$3–\$5 USD per click".
- 4 run live checkout; Copilot "does not take a commission or affiliate fee" (2026-09).

How often ChatGPT shows an ad: three measurements, unreconciled




Different metrics and markets; shown side by side, not averaged.

WHO BUYS

8 of 27 buyer segments show spend under either reading; willingness to pay is unknown in all 27

Segments by read, two counting rules (count of segments, of 27)



Filled = strict; outlined = loose. Measured-by-us, 2026-09-23.

How the 27 are built

Each segment crosses a vertical, a sub-market and a buyer size.


Two counting rules

Strict counts "none" only when every signal ran. *Loose* also counts a blocked, logged channel as checked. An earlier count exists (appendix).

Spend is present, but no source shows what any buyer paid.


Spend concentrates in organic tools, in large enterprises, and is read mostly from job postings

Spend segments by sub-market (count)



Measured-by-us, 2026-09-23

Local and multi-location, 9 further segments (count)



Measured-by-us, 2026-09-23

Of 7 spend reads counted 2026-09-22

- 6 rest on job postings.
- 5 are enterprise.
- 1 of 9 small and medium business (SMB) segments read spend.

The same postings, two readings

Segment files read them as spend; source censuses as attention. Both stand.

Buyer attention is high: "the No. 1 area of increased focus among buyers at 76%", of "more than 200" decision-makers (Interactive Advertising Bureau, 2026-09-10).

DOES IT WORK

No causal case exists, the correlational cases shrink under strict grading, and no brand confirms a vendor's case

Gold (causal) cases

0

in ~980 screened (2026-09-22) and ~1,640 more (2026-09-23)

Silver (before/after with control): two counts stand

7 or 1

7 as graded at collection; 1 when a case missing any of the 7 evidence-bar items is Bronze at best

Silver cases running negative or null

2 of 7

(filed; preprint, pulled 2026-09-22)

Brands confirming vendors' published cases

0 of 59

checked; all 59 silent

Published cases are survivors. Full-page re-grading moved 12 up and 0 down.

DOES IT WORK: WHAT THE EVIDENCE IS MADE OF

Most evidence is observational, brands that moved changed content rather than buying ads, and a filed revenue figure runs negative

52 cases by evidence class (count of cases)



Control cases counted observational; see appendix

What 108 brands that moved changed (count of brands)



One case can name several changes

Filed negative

NerdWallet credit-card revenue "decreased 24%" (filed, 2026-02-25).

Vendor conversion claims (traffic quality, not sales lift)

- Adobe: "AI Conversion Now 42% Higher" than non-AI visits (vendor-reported, 2026-03).
- Shopify: "nearly 50% higher rates than organic search" (vendor-reported, first quarter 2026).

The best-evidenced risk is that answers replace clicks; no source states likelihood, impact or mitigation for any risk

Rank	Risk	Strongest evidence
1	Referral collapse: answers replace clicks	"83-93% drops in click-through rates" (filed, 2026-09-17)
2	Regulation	EU AI Act Article 50, 2026-08-02; UK fair-ranking requirement, 2026-06-17 (filed)
3	Vendor consolidation	Adobe–Semrush 2026-04-28; Yext–GoShine 2026-09-01 (filed)
4	Platform capture: engines sell ads, checkout, measurement	OpenAI: "more than 50 technology and measurement partners" (company-stated)
5	No payback evidence	"just 3%" of \$250K spent (vendor-reported, 2026-03)
6	Measurement opacity	"0 of 34 rostered vendors disclose a fixed published prompt set" (vendor-reported)
7	"Do nothing" suffices	Google: "no additional requirements... nor other special optimizations necessary" (company-stated)
8	Engine refuses or withdraws inventory	Claude ad-free, 2026-02-04 (company-stated); Perplexity "winding down" ads (relayed, tier 5, 2026-02-18)

Ranked by strength of evidence that the risk is real, then breadth.

The strongest counter-arguments each rest on engine or vendor statements, not on brand outcomes

Question	What the evidence says
ChatGPT ads earn a "\$1 billion" run rate. Is that not proof?	It proves engines sell ads. It does not show brands gain: 0 causal cases; 0 of 59 brands confirm vendor cases.
Investors valued Profound at \$1.8B. Is demand not obvious?	Valuation is vendor-reported. Filed revenue exists for 2 of 34 vendors, neither product-specific.
AI visitors convert 42% better. Is that not a sales effect?	Vendor-reported traffic quality, not sales lift; no Gold case behind it.
Will forecasts settle the size question?	15 of 18 size figures are forecasts; agentic forecasts spread 26.3× (or "roughly 35×").
Could a brand simply do nothing?	Google states "no additional requirements... nor other special optimizations necessary". Ranked risk 7.

Prices paid, rate cards, measured size and sales effect are unknown; 11 of 25 load-bearing claims reach tier 3 or stronger

Unknown

- Price any buyer paid: unknown — not in repo.
- Engine rate cards: unknown — not in repo.
- Any measured sub-market size: unknown — not in repo.
- Agent-checkout protocol fee clauses: 6 of 6 unknown — checked 2026-09-22.

Gaps in our own checks

- Our own sampling of assistants stopped at a partial first day, 2026-09-22 22:40; ChatGPT 0 runs.
- Five data channels blocked, not exhausted, 2026-09-22.

Load-bearing claims at tier 3 or stronger

11 of 25 = 44.0%

(measured-by-us, 2026-09-23; an earlier count is in the appendix). The 80% bar is met on none of six counts.

Caveats

- Prices, counts, funding and cases are vendor or company statements; none audited.
- Published cases are survivors.

No approval is requested; three readings are yours to choose, and each is shown both ways

Reading you own	Option A	Option B
Which Silver count governs	7, as graded at collection	1, under strict evidence-bar grading
Which segment counting rule	Strict: 8 spend · 1 attention · 17 none · 1 blank	Loose: 8 spend · 1 attention · 18 none · 0 blank
Which sub-market a hypothetical product would address	Organic: 6 spend segments · Paid: 1 · Agentic: 1. Organic has disclosed floors; paid has engine revenue; agentic has neither.	

The evidence supports: engines sell placement, and some enterprises pay for organic tools. It stops before: what buyers paid, how big any sub-market is, and whether visibility raises sales.

Where one figure has two readings or two dates, both stand

Figure	Earlier or alternative reading	Later or main reading
27-segment tally	7 spend · 1 attention · 11 none · 8 blank, 2026-09-22	strict 8·1·17·1; loose 8·1·18·0, 2026-09-23
Organic spend segments	5, 2026-09-22	6, 2026-09-23
52 cases by class (experimental · observational · action only)	10 · 29 · 13 (control counted experimental), 2026-09-23	5 · 34 · 13 (control counted observational), 2026-09-23
108 movers by class (experimental · observational · action only · outside)	6 · 61 · 37 · 4 (control counted experimental), 2026-09-23	1 · 66 · 37 · 4 (control counted observational), 2026-09-23
Agentic forecast spread	"roughly 35×", 2029–2030	26.3×, 2030
Skincare agentic, enterprise segment	spend (segment file)	attention (sub-market file)
Claims at tier 3 or stronger	17 of 25 = 68.0%, 2026-09-22	11 of 25 = 44.0%, 2026-09-23

Sources: measured-by-us, 2026-09-22 and 2026-09-23; analyst-derived forecasts, pulled 2026-09-22

Terms and abbreviations

- **AI**: artificial intelligence. **EU / UK / US**: European Union, United Kingdom, United States.
- **AI Act Article 50**: EU transparency duties for AI systems.
- **Very large online search engine**: a designation under the EU Digital Services Act.
- **ARR**: annual recurring revenue. **Run rate**: current revenue annualised.
- **SMB**: small and medium business.
- **USD**: US dollars. **M / B / T**: million / billion / trillion.
- **Rate card**: a published price list for ad placements.
- **IAB**: Interactive Advertising Bureau. **EMARKETER, WPP Media, Similarweb, Adthena**: research, media and panel firms.
- **Series D**: a late-stage venture funding round.
- **Rostered vendor**: a tool vendor on the programme's tracked list.
- **Load-bearing claim**: a claim the conclusions depend on.
- **Evidence bar**: the checklist a case must meet to be graded Silver or Gold.
- **Relayed**: a figure reported by press, not confirmed by the parties.