

EVIDENCE BRIEFING FOR A DIRECTOR

Brands inside AI assistants: is anyone paying to be recommended, and does it work?

Demand for brand visibility and product recommendation inside ChatGPT, Gemini, Perplexity, Copilot and Claude, read segment by segment from published evidence.

Research only: evidence, not a verdict

Evidence as of 2026-09

The decision belongs to the reader. Internet sources only; no interviews. Acronyms are defined on first use and in the glossary at the back.

Brands have started paying to appear inside AI assistants, but no published evidence shows it pays back

Situation. AI (artificial intelligence) assistants reach audiences above one billion users each.

Complication. In 2026 they began selling placement inside answers, while publishers lose clicks to those answers.

Question. Which segments pay to be seen and recommended inside AI assistants, and does anything show it works?

What is asked of you. No decision is requested. Three reading choices move the figures; they are set out near the end.

The evidence supports three statements and stops at a fourth

1. Paid inventory exists and is scaling

ChatGPT ads: "\$1 billion in annualized revenue run rate" (company-stated, 2026-08-31). No engine publishes a rate card or payback.

2. Buyer spend is real but thin and early

Spend shows in 8 of 27 segment cells, strict reading (analyst-derived, 2026-09-23). No buyer discloses a price paid.

3. Tool vendors are funded and consolidating

Profound: "\$180M Series D at \$1.8B valuation" (vendor-reported, 2026-09-15). Disclosed product revenue is small.

4. It stops at proof

About 980 candidate cases screened; zero disclose a design that isolates cause (measured-by-us, 2026-09-22).

Every figure carries a source kind, every source a provenance tier, and every success case a grade

Source kind: what sort of number it is

filed	regulator, court or securities filing
company-stated	the company's own page, release or call
vendor-reported	a vendor or agency about its own market
analyst-derived	compiled or forecast by analysts, including this team
measured-by-us	counted by this research team from pulled sources

Provenance tier: lower is stronger

1	measured by us
2	filed: securities filing, court exhibit
3	platform primary: own docs, pricing page
4	panel or telemetry, method published
5	vendor study with sample size, dates, method
6	vendor blog, no sample: "Marketing. Not a source"
7	listicle or roundup: not pulled

Case grade: does it show cause?

Gold	holdout, geo-split or switchback; causal
Silver	before and after, plus an unaffected control; correlational
Bronze	visibility or mention change only; no revenue link
Fools gold	revenue claim, no baseline or no control

Holdout: a group kept unexposed. **Geo-split:** regions compared. **Switchback:** turned on and off in turns.

The market splits three ways, read separately and never added together

Organic recommendation

A brand is named or cited in an assistant's answer, unpaid.

Paid placement

Ads shown in or beside an assistant's answer.

Agentic commerce

The assistant completes the purchase for the user.

Demand was read in 27 cells: 3 sub-markets × 3 verticals × 3 buyer sizes

Vertical	What it covers
Skincare and beauty	consumer brands
B2B SaaS	business-to-business software sold as a subscription
High-CPA regulated	regulated categories with high cost per acquisition (CPA), such as insurance

Buyer sizes: SMB (small and medium-sized business), mid-market, enterprise. A fourth vertical, local and multi-location businesses, was read separately.

Each cell reads one of four ways

spend: money committed

attention: interest, no spend

none: checked, nothing found

blank: not yet checked

Assistants reach over a billion users each, now sell ads, and take clicks from publishers

Reach

> 1 billion


ChatGPT weekly active users
company-stated, 2026-08-31

1 billion

Google AI Mode "surpassed 1
billion monthly active users"
company-stated, 2026-05-19

Weekly and monthly users are different
metrics; shown apart, not compared.

Placement sold: ChatGPT ads, live 2026-02-09



Advertisers: "over 600" (2026-03-26) to "tens of thousands" (2026-08-31).

Clicks lost

83-93%

"drops in click-through rates" on
news domains, Microsoft data in a
court filing
filed, 2026-09-17

Click-through rate (CTR): share of viewers
who click a link.

2026 turned AI answers into an ad and checkout channel, and EU duties took effect the same year

Date	Type	Event	Source kind
2026-01-08	● Engine	Copilot Checkout live, US; "does not take a commission"	company-stated
2026-02-04	● Engine	Anthropic: "Claude will remain ad-free"	company-stated
2026-02-09	● Engine	ChatGPT ads live, US pilot	company-stated
2026-04-28	■ Deal	Adobe completes Semrush purchase, "\$12.00/share cash"	filed
2026-08-02	◆ Regulation	EU AI Act Article 50 transparency duties in force	filed
2026-08-18	● Engine	ChatGPT ads expand to 31 European markets	company-stated
2026-08-31	◆ Regulation	EU designates ChatGPT a VLOSE; duties from 2027-01	filed
2026-09-15	■ Deal	Profound raises "\$180M Series D at \$1.8B valuation"	vendor-reported

EU: European Union. US: United States. VLOSE: very large online search engine, a designation under the EU Digital Services Act that brings extra duties.

No sub-market has a measured size: only disclosed floors, and forecasts that disagree widely

Of 18 published size figures, **15 are forecasts and 0 are measured sizes**. They sit at provenance tier 5–6; tier 6 reads "Marketing. Not a source". None is used here as a size.

Organic recommendation

Disclosed floor, 2026-09

\$42.2M–\$48.2M + €2.2M annual recurring revenue (ARR) (vendor-reported, filed)

Alternative floor: **\$5.3M–\$35.2M**, list price × customer count (vendor-reported inputs)

Forecast for 2034

\$17.15B–\$32.92B, 4 forecasts (analyst-derived)

Paid placement

Disclosed floor, 2026-08

≥ \$1B annualized run rate, one engine (company-stated)

Forecast for 2030

"just over \$5 billion" to "over \$100 billion" (analyst-derived)

Agentic commerce

Disclosed floor

unknown — not in repo; 0 engines disclose a merchant count (company-stated, 2026-09)

Forecast for 2030

\$190B–\$5T, 7 forecasts (analyst-derived). Spread: 26.3× (2030 only) or "roughly 35×" (2029–30)

The two organic floors use different bases and are not reconciled. Floors and forecasts are different metrics and are not charted together.

Spend shows in 8 of 27 segment cells; 6 are organic, and only 1 is a small business

Vertical	Organic			Paid			Agentic		
	SMB	Mid	Ent.	SMB	Mid	Ent.	SMB	Mid	Ent.
Skincare and beauty	ATTN	blank	SPEND	none	none	SPEND	none	none	SPEND*
B2B SaaS	SPEND	SPEND	SPEND	none	none	none	none	none	none
High-CPA regulated	none	SPEND	SPEND	none	none	none	none	none	none

Strict reading, 2026-09-23. SPEND = money committed; ATTN = attention only; none = checked, nothing found; blank = not checked (dashed). Mid = mid-market; Ent. = enterprise.

* Skincare agentic enterprise reads spend here and attention in the agentic-market read; both stand.

Two readings of empty cells, 27 cells

Reading	spend	attn.	none	blank
Strict	8	1	17	1
Loose	8	1	18	0

Strict: a cell reads none only if every signal ran. **Loose:** a channel blocked and logged counts as checked. The one strict blank, skincare organic mid-market, reads none under loose.

Spend by sub-market: organic 6, paid 1, agentic 1.

The spend evidence is mostly hiring at large firms, and no buyer discloses a price paid

First read of spend cells, 2026-09-22 (7 spend cells then)

6 of 7

spend reads rested on
job postings

5 of 7

spend cells were
enterprise

1 of 9

SMB cells read spend

Job postings are the weakest spend class: "headcount budget, not category spend". A spend cell shows money moved in it, never how much.

Willingness to pay: unknown in all 27 cells (2026-09-23). No buyer in any vertical discloses a price paid.

Local and multi-location businesses, 9 cells read outside the 27, 2026-09-23

Read	Cells
spend	1
attention	3
none, checked	5
blank	0

The one local spend cell is enterprise organic, again a job posting. Strict and loose readings agree here.

Engines own the ad and checkout inventory, yet publish almost no prices or revenue

Of 8 engines, as of 2026-09 (● yes, ○ no)

Live ad product ● ● ● ● ● ○ ○ ○ 5 of 8

Publishes a rate card ○ ○ ○ ○ ○ ○ ○ ○ 0 of 8

Discloses ad revenue ● ○ ○ ○ ○ ○ ○ ○ 1 of 8

Of 4 live assistant checkout programs, as of 2026-09

States a merchant fee ● ○ ○ ○ 1 of 4

The only price stated

OpenAI bid guidance: "\$3–\$5 USD per click" (company-stated, read 2026-09-22)

Checkout without a fee



Copilot Checkout "does not take a commission" (company-stated, 2026-01-08)

One engine stays out

Anthropic: "Claude will remain ad-free" (company-stated, 2026-02-04)

Tool vendors are funded and consolidating, but few disclose prices and almost none revenue

Of 34 rostered visibility-tool vendors, 2026-09



Bar length: vendor count, linear from 0; grey = all 34

A prompt set is the list of questions a tool asks assistants to measure visibility.

Money and deals

Event	Source, date
Profound: "\$180M Series D at \$1.8B valuation"	vendor-reported, 2026-09-15
Adobe completes Semrush purchase, "\$12.00/share cash"	filed, 2026-04-28
Sitecore buys Scrunch; value not disclosed, \$225M per press, unconfirmed	company-stated, 2026-06-03
Semrush AI products "surpassed \$38 million in ARR"	filed, 2026-03-02
Peec AI ARR \$4M (2025-11), then \$10M (2026-05)	vendor-reported

No published case isolates a sales effect: none of about 980 candidates disclosed a causal design

~980

candidate cases screened; **zero** disclose a holdout, geo-split or switchback
measured-by-us, 2026-09-22


Gold: 0

Silver: 7 or 1

7 as first graded; 1 under the literal evidence bar
measured-by-us, 2026-09

52 classed cases, by evidence class: two readings, 2026-09-23

Reading 1



Reading 2



Bar length: case count, linear from 0

■ experimental design ■ observational ▨ action, outcome unknown

Reading 1 counts a before/after with a control as experimental; **Reading 2** counts it as observational. Neither makes it causal: such cases stay Silver.

Negative and weak results sit beside the positives, and none of the positives is causal

Negative, null or weak


NerdWallet credit cards revenue "decreased 24%"	filed, 2026-02-25
Replication: optimisation levers move citation on "none" of ten engine families	preprint, analyst-derived, 2026-09-07
SE Ranking's own ChatGPT ad campaigns: CTR 1.30% on 97,000+ impressions	vendor-reported, 2026-08-10
One ChatGPT advertiser spent "just 3%" of a \$250K budget	vendor-reported relay, 2026-03
Ad results with a control: 0 of 15	analyst-derived, 2026-09-23

Positive, none causal

Adobe: "AI Conversion Now 42% Higher" than non-AI visits	vendor-reported, 2026-03
TechTarget: "2x to 3x higher membership conversion rate"	filed, 2026-03-11
Both compare AI-referred visits with other visits at one point in time: correlation, not a measured lift.	

Brands that moved changed content and tools; only 2 of 108 cases name paid media

Of 108 cases, 99 name a change. Cases naming each kind, 2026-09-23



Bar length: number of cases, linear from 0.

A case can name more than one kind; bars do not add to 99.

Content: new or rewritten pages, structured data. **Vendor tool:** a tool adopted. **Organisation:** hire or team. **Other:** agency, PR, targets, checkout.

Both paid-media cases are one Google pilot.

Another count of the same 108 exists: 67 metric moved, 37 action only, 4 neither.


Eleven ranked risks; lost clicks and regulation lead, and none carries a likelihood

Rank	Risk	Strongest evidence
1	Referral collapse: answers replace clicks	"83-93% drops in click-through rates" (filed, 2026-09)
2	Regulation	EU AI Act Article 50 in force 2026-08-02; UK fair-ranking order 2026-06-17 (filed)
3	Vendor consolidation	Adobe–Semrush 2026-04-28 (filed); Sitecore–Scrunch 2026-06-03 (company-stated)
4	Platform capture: engines sell ads, checkout, measurement	OpenAI: "more than 50 technology and measurement partners" (company-stated, 2026-08-31)
5	No payback evidence	0 Gold; 0 of 15 ad results controlled (analyst-derived, 2026-09-23)
6	Measurement opacity	0 of 34 vendors publish a fixed prompt set (vendor-reported, 2026-09)
7	Substitute: do nothing	Google: "no additional requirements... nor other special optimizations necessary"

Ranked by the research team, 2026-09-23 (analyst-derived): strongest evidence first, then breadth of cells touched. Likelihood, impact, mitigation and owner: no source states one. Ranks 8–11 in the appendix. UK: United Kingdom.

Claim strength falls short of the programme's own 80 percent bar under every count

Load-bearing claims at provenance tier 3 or better (measured, filed or platform primary)



Share of claims, linear from 0.
Two lists on two tier bases; all counts stand.

Gaps that limit every reading

- Own sampling of assistant answers held 2026-09-22 22:40 and skipped 2026-09-23; ChatGPT and Perplexity got 0 runs.
- Channels still blocked, 2026-09-23: 5; 3 more need payment. Absence is only as strong as these channels.
- Whether any posting was filled or change shipped: unknown — not in repo.

No decision is requested; three reading choices are yours, and each moves the figures

Choice	Option A	Option B	What it moves
How to read empty cells	Strict: 8 spend · 1 attention · 17 none · 1 blank	Loose: 8 spend · 1 attention · 18 none · 0 blank	Whether the demand map is complete; spend is unchanged
How many Silver cases	7, as first graded	1, under the literal evidence bar	How much correlational support exists
Evidence class, 52 cases	Reading 1: 10 experimental design	Reading 2: 5 experimental design	How much of the base looks designed; none is causal

Any product in this space is a hypothesis. The evidence does not pick its sub-market. The verdict on whether demand is sufficient sits with you.

Caveats: self-reported figures, survivor cases, and demand read from public traces only

- Prices, customer counts, funding and case results are self-reported; none is audited.
- Published cases are survivors; quiet failures leave no trace. Full-page re-grading moved 12 cases up and 0 down (measured-by-us, 2026-09).
- A spend cell shows money moved in it, never how much. Quiet spending leaves no public trace, so none and blank are weaker evidence than spend.
- Tiers rank provenance, not correctness: a filing still carries the filer's framing.
- Internet sources only: no interviews, no outreach. Evidence frozen at 2026-09-23; nothing later is used.
- Figures with two readings or two dates are paired in the appendix; neither is dropped.

Figures carried two ways (1 of 2): both readings kept, neither averaged

Figure	First reading	Second reading
Cell reads, 27 cells	7 spend · 1 attention · 11 none · 8 blank (2026-09-22)	Strict 8/1/17/1; loose 8/1/18/0 (2026-09-23)
Spend cells by sub-market	Organic 5 · paid 1 · agentic 1 (2026-09-22)	Organic 6 · paid 1 · agentic 1 (2026-09-23)
Silver count	7, as graded on first read	1, under the literal evidence bar (2026-09-22)
Evidence class, 52 cases	Reading 1: 10 / 29 / 13 (2026-09-23)	Reading 2: 5 / 34 / 13 (2026-09-23)
108 moved cases	99 name a change	67 metric moved · 37 action only · 4 neither (2026-09-23)
Organic floor	"\$42.2M–\$48.2M + €2.2M" disclosed ARR	"\$5.3M–\$35.2M" list price × customer count
Agentic forecast spread	26.3× (2030 only)	"roughly 35×" (2029–30)

Figures carried two ways (2 of 2): later figures do not erase earlier ones

Figure	First reading	Second reading
ChatGPT advertisers	"over 600" (2026-03-26, press relay)	"tens of thousands" (2026-08-31, company-stated)
ChatGPT ads revenue	"\$100 million annualized" (2026-03-26, press relay)	"\$1 billion" run rate (2026-08-31, company-stated)
Peec AI ARR	\$4M (2025-11, vendor-reported)	\$10M (2026-05, vendor-reported)
Claims at tier 3 or better	17 of 25 = 68.0% (2026-09-22)	12 of 27 = 44.4%; 11 of 25 = 44.0% (2026-09-23)
Scrunch deal value	not disclosed (Sitecore)	\$225M (Bloomberg, unconfirmed)
Skincare agentic × enterprise	spend (customer read)	attention (agentic-market read)

Risks ranked 8 to 11: engine reversals, checkout fees, manipulation, and an unsizeable market

Rank	Risk	Strongest evidence
8	Engine policy reversal: inventory withdrawn or refused	"Claude will remain ad-free" (company-stated, 2026-02-04); Perplexity ads "winding down" (press relay, pulled 2026-09-23)
9	Agentic rails: checkout on the engine, fees undisclosed	6 of 6 protocol fee clauses unknown (company-stated, 2026-09)
10	Manipulation and countermeasures	Google spam policy names "manipulate generative AI responses", 2026-08-28 (company-stated)
11	Unsizeable market: every published size is a forecast	18 figures, 15 forecasts, 0 measured sizes (tier 6 table, 2026-09-22)

Same ranking rule as ranks 1–7 (analyst-derived, 2026-09-23). Likelihood, impact, mitigation and owner: no source states one.

Glossary: acronyms and reading labels used in this deck

AI: artificial intelligence. **AI Mode, AI Overviews:** Google's AI answer surfaces.

ARR: annual recurring revenue.

Article 50: EU AI Act article setting transparency duties.

B2B SaaS: business-to-business software sold as a subscription.

CPA / High-CPA: cost per acquisition; categories where each customer costs a lot to win.

CTR: click-through rate.

EU, UK, US: European Union, United Kingdom, United States. **USD:** US dollars.

Rate card: a published price list for ads.

Run rate: current revenue pace, stated per year.

SMB: small and medium-sized business.

VLOSE: very large online search engine, an EU Digital Services Act designation.

Cell read: spend (money committed), attention (interest, no spend), none (checked, nothing found), blank (not checked).

Strict / loose: strict reads none only if every signal ran; loose also counts blocked-and-logged channels as checked.

Reading 1 / Reading 2: a before/after with control counted as experimental, or as observational.

Gold / Silver / Bronze / Fools gold: case grades, defined on slide 3.

Preprint: a research paper not yet peer-reviewed.

Series D: a late-stage venture funding round.